Response Style Variation Across Model Families: Evidence from Large-Scale Psychometric Probing of LLMs

Anonymous Authors

Abstract

A growing body of work applies personality inventories to LLMs, yet it remains unclear whether the resulting scores capture meaningful individual differences or systematic response tendencies shaped by training and alignment. In this study, we administer 61 Likert-scale items spanning 9 psychometric dimensions to one model from each of 15 model families via unified APIs, collecting 10,980 responses (12 seeds per model), and characterize the resulting response style variation. We find that model families exhibit large, systematic differences in their response distributions (median pairwise Cohen's $d = 1.08$), with HEXACO-H showing the largest between-family effect ($\eta_p^2 = 0.83$). The inter-dimension correlation structure reorganizes relative to humans: the Conscientiousness–Neuroticism correlation reverses from $r = -0.25$ to $r = +0.77$ (95% CI $[0.42, 0.92]$), and Extraversion–Openness flips from positive to negative. Convergent validity against published human baselines yields insufficient evidence for construct equivalence (3/5 sign agreement, $p = 0.50$), further indicating that standard dimension labels do not transfer cleanly to LLMs. Together, these findings demonstrate systematic response style variation across LLM families and quantify the degree to which psychometric responses are shaped by training and alignment rather than by human-like individual differences.

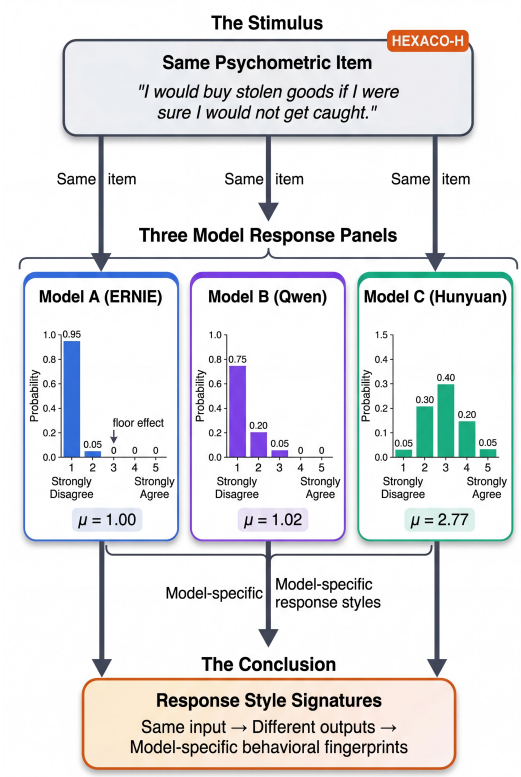

Figure 1: **Core concept.** Identical psychometric items elicit systematically different Likert-scale responses from different models. These model-specific response style signatures raise the question: do psychometric scores in LLMs reflect personality traits, or systematic response tendencies shaped by training and alignment?

1 Introduction

A growing number of studies apply personality inventories to LLMs, reporting that different models exhibit different trait profiles (Jiang and Althoff, 2023; Pan and Zeng, 2024; Serapio-García et al., 2025). These results are often interpreted as evidence that LLMs possess personality-like traits shaped by training and alignment. However, if LLMs respond to personality items using surface-level heuristics—such as mimicking socially desirable answers or following alignment-induced response tendencies—then the resulting scores may capture systematic *response styles* rather than personality traits, and the inter-dimension correlations may differ fundamentally from the human factor structure. Following Binz and Schulz (2023), we use the term *response style* to describe observed response patterns without claiming they reflect personality traits or inner states.

We characterize this response style variation with a large-scale cross-model comparison (Figure 1). Using unified API platforms, we administer 61 Likert-scale items spanning 9 dimensions—the Big Five Inventory (John and Srivastava, 1999), HEXACO Honesty-Humility (Lee and Ashton, 2004), Schwartz Values (Schwartz, 1992), cognitive style, and cultural dimensions—to one model from each of 15 model families, repeating each query across 12 random seeds (10,980 total API calls). We also track 16 additional models from three families (Qwen, DeepSeek, GLM) across versions, yielding 31 models total.

We present three findings that characterize systematic response style variation across LLM families:

1. **Model families exhibit distinct response style signatures.** Families differ substantially in their response distributions (median pairwise Cohen's $d = 1.08$), indicating that different training approaches produce systematic behavioral fingerprints.
2. **Inter-dimension correlations reorganize relative to humans.** The Conscientiousness–Neuroticism correlation reverses from $r = -0.25$ in humans to $r = +0.77$ in LLMs (95% CI $[0.42, 0.92]$), and Extraversion–Openness flips from positive to negative, patterns observed across most tested subgroups.
3. **Standard dimension labels do not transfer cleanly.** Comparison against published human baselines yields only 3/5 sign agreement ($p = 0.50$); this inconclusive convergent validity result itself constitutes evidence for trait structure reorganization rather than random noise.

We test convergent validity against human baselines (§4.5), correct for acquiescence bias (§4.7), and report 95% confidence intervals for key statistics.

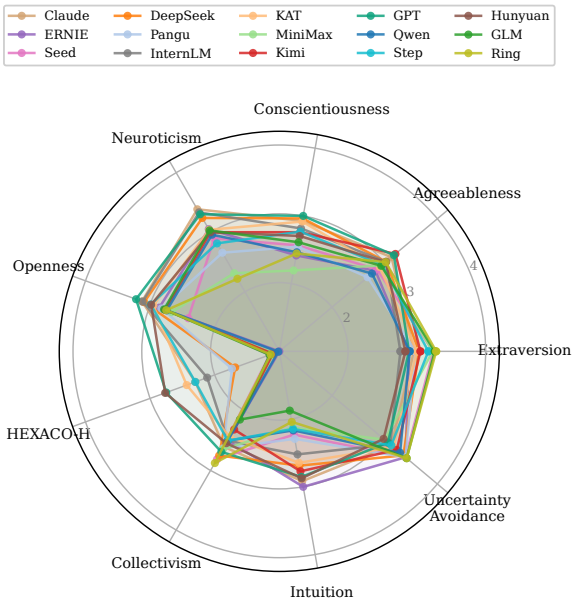

Figure 2: Response style profiles of all 15 model families across 9 dimensions. Each line is one model family. The distinct shapes suggest that different training approaches leave different behavioral fingerprints on models.

2 Related Work

LLM personality and response behavior. Jiang and Althoff (2023) found temperature-dependent personality variation in ChatGPT using the BFI-10. Pan and Zeng (2024) compared LLMs from different cultural backgrounds on the BFI-44, reporting cultural effects. Myska and Macková (2023) showed that GPT-3.5 and GPT-4 produce stable profiles across sessions, and Schaeckermann et al. (2024) examined 9 LLMs on the Big Five, finding moderate-to-large inter-model differences. Serapio-García et al. (2025) proposed a psychometric framework for evaluating and shaping personality traits across 18 LLMs, reporting that larger models exhibit broader trait ranges. Hilliard et al. (2024) confirmed that model scale affects elicited personality profiles, with larger models showing more extreme trait scores. Bodroža et al. (2024) tested 7 LLMs on multiple personality instruments, finding limited temporal stability but consistently elevated prosociality. Santos and Costa (2024) surveyed personality assessment with LLMs, highlighting both promises and pitfalls of treating model outputs as psychometric data. Sorokovikova et al. (2024) replicated Big Five profiles across three models and found systematic deviations from human norms, while Tosato et al. (2025) documented persistent instability across scales and conversation histories.

While Serapio-García et al. (2025) evaluated the largest set of models to date (18), their study used heterogeneous APIs and did not test for structural differences in inter-dimension correlations. Our work contributes unified API control across 15 model families and explicit validity testing against human baselines.

Behavioral evaluation beyond capabilities. Standard capability benchmarks (Hendrycks et al., 2021; Liang et al., 2022) measure task performance—reasoning, coding, instruction-following—and instruction-following evaluators such as AlpacaEval (Dubois et al., 2024) assess response quality, but none address how models respond to self-referential statements about preferences and moral tendencies. Laverghetta et al. (2022) showed that language models can predict human psychometric properties from text, raising the question of whether LLM self-reports reflect similar constructs. We contribute a behavioral evaluation dimension: how models respond to standardized self-referential questions.

Response biases and measurement validity. Applying human-designed psychometric instruments to LLMs raises validity concerns (Schaeckermann et al., 2024; Serpell and Ney, 2024). Sühr et al. (2024) question whether standard personality tests are valid for LLMs at all, given the lack of genuine self-report. Salecha and Kosinski (2024) demonstrated that LLMs exhibit human-like social desirability biases in Big Five survey responses, inflating socially desirable traits. Domínguez-Olmedo et al. (2024) showed that LLM survey responses lack the natural variation seen in human populations, producing compressed response distributions. Tjuatja et al. (2023) found that LLMs fail to replicate known human response biases in survey design. Braun (2025) documented acquiescence bias in LLMs, systematically agreeing with questionnaire statements regardless of content. Han et al. (2025) demonstrated a dissociation between LLM self-reported personality and behavioral choices, reinforcing the need to distinguish stated preferences from actions. Xie et al. (2025) introduced AIPsychoBench, a benchmark specifically designed to measure psychometric differences between LLMs and humans, finding systematic discrepancies in response patterns. We address these concerns by framing measurement explicitly as "response style" rather than personality, testing convergent validity against human data, and correcting for acquiescence bias—extending methodological transparency beyond prior work.

Table 1: Psychometric instruments: 9 dimensions, 61 Likert items (1–5 scale).

Dimension	Source	Items	Reverse Scored
Extraversion	BFI-44	8	4 items
Agreeableness	BFI-44	9	4 items
Conscientiousness	BFI-44	9	4 items
Neuroticism	BFI-44	8	2 items
Openness	BFI-44	10	3 items
HEXACO-H	HEXACO	5	0
Collectivism	Schwartz	4	0
Intuition	Cognitive Style	4	2 items
Uncertainty Avoidance	Cultural Dimensions	4	2 items
Total		**61**	

Factor structure in LLMs. The Five-Factor Model (McCrae and Costa, 1992) provides the theoretical foundation for the BFI dimensions we use. Suh and Moon (2024) recovered the Big Five factor structure from LLMs using singular value decomposition on log-probabilities of personality-relevant adjectives. Their approach analyzes semantic space, whereas we analyze questionnaire responses—two complementary probes. An exploratory PCA on our data hints at a three-factor reorganization (Appendix F), but this is severely underpowered ($n = 15$).

3 Method

3.1 Psychometric Instruments

We use 9 dimensions comprising 61 Likert items (1–5 scale; Table 1). Five come from the Big Five Inventory (John and Srivastava, 1999; McCrae and Costa, 1992; O'Connor, 2007): Extraversion (8 items), Agreeableness (9), Conscientiousness (9), Neuroticism (8), and Openness (10). We add HEXACO Honesty-Humility (5 items; Lee and Ashton, 2004), Collectivism (4 items; Schwartz Values Survey Schwartz, 1992), Intuition (4 items; cognitive style), and Uncertainty Avoidance (4 items; cultural dimensions). Reverse-scored items are transformed via score $= 6 -$ original; dimension scores are the mean of constituent items.

3.2 Model Selection

We select 15 model families accessible through unified API platforms, contributing one representative model each (Table 2). Priorities: flagship models where available; Dense and MoE architectures both represented (4 Dense, 9 MoE, 2 undisclosed).

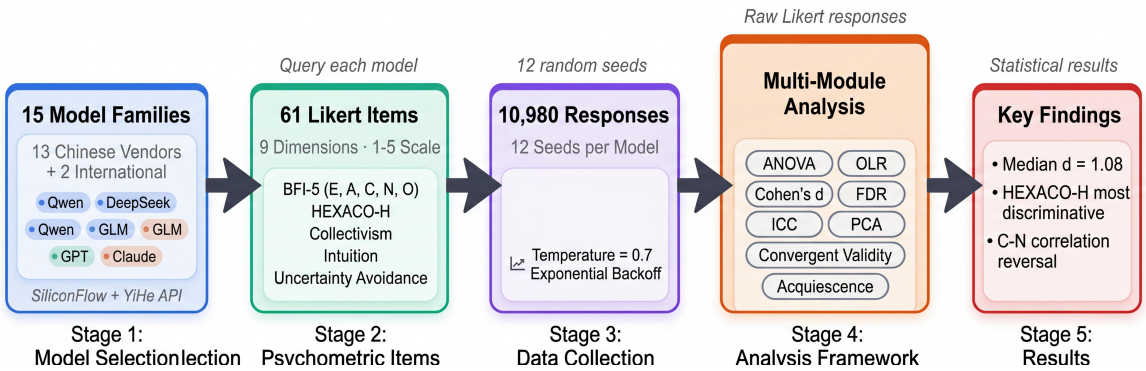

Figure 3: **Experimental pipeline.** We select 15 model families via unified API platforms, query each model on 61 Likert items (9 dimensions, 12 seeds), collect 10,980 responses, and apply a multi-module analysis framework (ANOVA, Cohen's d, convergent validity, reliability, acquiescence correction) to characterize model-level response style variation.

Table 2: Models from 15 model families, one representative model per family. Total parameters in parentheses where publicly disclosed.

#	Model Family	Model ID	Architecture
1	Qwen	Qwen3.5-397B-A17B	MoE (397B)
2	DeepSeek	DeepSeek-V3.2	MoE (671B)
3	GLM	GLM-5	MoE (744B)
4	Kimi	Kimi-K2.5	MoE (1.1T)
5	ERNIE	ERNIE-4.5-300B-A47B	MoE (300B)
6	Hunyuan	Hunyuan-A13B-Instruct	Dense (13B)
7	Seed	Seed-OSS-36B-Instruct	Dense (36B)
8	InternLM	internlm2_5-7b-chat	Dense (7B)
9	Ring	Ring-flash-2.0	MoE (100B)
10	Step	Step-3.5-Flash	MoE (197B)
11	Pangu	pangu-pro-moe	MoE (72B)
12	KAT	KAT-Dev	Dense (32B)
13	MiniMax	MiniMax-M2.5	MoE (230B)
14	GPT	GPT-5	Undisclosed
15	Claude	Claude Opus 4.5	Undisclosed

3.3 Experimental Protocol

Each model is administered all 61 items per seed at temperature $= 0.7$, with 12 seeds (0, 1, 2, 4, 5, 6, 7, 8, 9, 42, 123, 456) capturing response distribution variation, totaling 10,980 API calls. We set temperature to 0.7 to balance response diversity with reproducibility: temperature 0 yields deterministic outputs that eliminate within-model variance, while higher temperatures (≥ 1.0) may produce inconsistent self-referential responses. Prior work on LLM personality (Jiang and Althoff, 2023) has shown that temperature affects reported trait scores, making a moderate temperature necessary to capture the full range of each model's response tendencies. Items are framed as self-referential statements (*e.g.* "I see myself as someone who is the life of the party") with a 1–5 agreement scale. Rate limiting uses exponential backoff on 429 responses with up to 5 concurrent threads.

3.4 Analysis Framework

Figure 3 illustrates the full experimental pipeline. Our analysis combines several complementary approaches. For each dimension, we run a one-way ANOVA (Model Family, 15 levels) with FDR-corrected p-values (Benjamini and Hochberg, 1995) and partial η_p^2, supplemented by ordered logistic regression to account for the ordinal nature of Likert responses. The effect size is quantified as partial η_p^2:

$$\eta_p^2 = \frac{SS_{\text{between}}}{SS_{\text{between}} + SS_{\text{within}}} \tag{1}$$

where SS_{between} and SS_{within} are the between-group and within-group sums of squares, respectively. Pairwise Cohen's d (Welch's t-test) is computed for all 105 model pairs across 9 dimensions (945 tests, FDR-corrected):

$$d_{ij} = \frac{\bar{x}_i - \bar{x}_j}{s_{\text{pooled}}}, \quad s_{\text{pooled}} = \sqrt{\frac{(n_i - 1)s_i^2 + (n_j - 1)s_j^2}{n_i + n_j - 2}} \tag{2}$$

where $\bar{x}_i$ and $\bar{x}_j$ are the mean dimension scores for model families i and j, s_i and s_j are their standard deviations, n_i and n_j are the number of observations (12 seeds each), and s_{pooled} is the pooled standard deviation. Power analysis confirms adequate power ($> 80\%$) for $d \geq 0.8$ with 12 observations per model. To correct for acquiescence bias, we compute partial correlations controlling for the first principal component (PC1):

$$r_{XY \cdot Z} = \frac{r_{XY} - r_{XZ} r_{YZ}}{\sqrt{(1 - r_{XZ}^2)(1 - r_{YZ}^2)}} \tag{3}$$

where r_{XY} is the Pearson correlation between dimensions X and Y, r_{XZ} and r_{YZ} are their respective correlations with PC1 (Z), and $r_{XY \cdot Z}$ is the partial correlation between X and Y after removing the variance explained by PC1. We assess reliability via ICC(1,1) and coefficient of variation, classify dimensions with inter-model $\sigma < 0.15$ as alignment artifacts, and compare LLM inter-dimension correlations against published human baselines (John and Srivastava, 1999) for 5 key dimension pairs. For all correlation and ANOVA analyses, we first aggregate the 12 seed-level responses per model into dimension means, then compute statistics at the model-family level ($n = 15$). The human baseline correlations are drawn from the large-scale validation of the BFI-44 reported in the BFI-2 manual (O'Connor, 2007) and the trait taxonomy chapter (John and Srivastava, 1999); exact standard deviations for these inter-dimension correlations are not always available in the original sources.

4 Experiments

4.1 Overview

Table 3 reports model family mean scores across 9 dimensions. On the Big Five, scores cluster around the scale midpoint (2.7–3.3) with the exception of Neuroticism, where Ring (2.22) and MiniMax (2.31) stand apart from GPT (3.30) and Claude (3.39). Intuition shows the widest spread among non-H dimensions: GLM scores lowest (1.88) while Claude leads (2.92), a gap of over one full point. Collectivism and Uncertainty Avoidance are the most compressed dimensions (ranges of 0.73 and 0.44, respectively), suggesting near-uniform responses across model families. In the raw data, Dense models (Hunyuan, Seed, InternLM, KAT) tend to score higher on HEXACO-H and Agreeableness than MoE models, while MoE models show higher Neuroticism and Collectivism—a pattern we revisit in §4.4 with appropriate caution about the small Dense sample ($n = 4$). Figure 2 overlays all 15 profiles, revealing that models from the same provider tend to share similar profile shapes. Family-level clustering is further illustrated by within-family trajectories across model versions (Appendix G).

4.2 Cross-Model Discrimination

All 9 dimensions show significant model effects ($p < 0.001$, FDR-corrected; Table 5), with the strongest discrimination on HEXACO-H ($F = 56.75$, $\eta_p^2 = 0.83$) and Neuroticism ($F = 34.50$, $\eta_p^2 = 0.75$); non-parametric Kruskal-Wallis tests confirm these results. Figure 4 shows the inter-model σ by dimension: HEXACO-H has the highest variability (0.68), followed by Neuroticism (0.38) and Intuition (0.35), while only Agreeableness ($\sigma = 0.14$) falls below our alignment-artifact threshold.

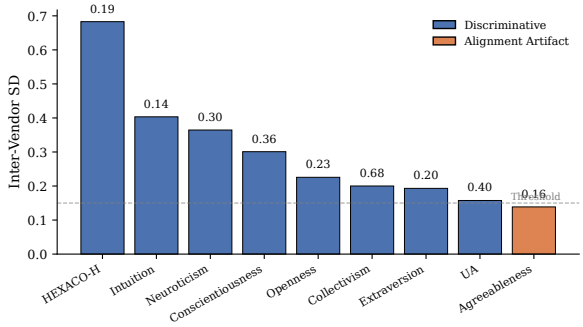

Figure 4: Inter-model SD by dimension. Blue: discriminative ($\sigma \geq 0.15$); orange: alignment artifact.

Architecture type (Dense vs. MoE) shows suggestive but non-significant differences (Table 4), with only 4 Dense families rendering all comparisons severely underpowered. The strongest signal is on HEXACO-H ($d = 1.44$, 95% CI $[0.14, 2.74]$, $p = 0.11$), but this is a descriptive pattern rather than a statistically supported conclusion. These patterns are explored further in §4.4, where we emphasize their exploratory nature.

The pairwise effect sizes are large: the median Cohen's d across all 105 model pairs and 9 dimensions is 1.08, and every pair exceeds $d \geq 0.8$ on at least one dimension (Figure 5). Table 6 reports the largest pairwise Cohen's d for each dimension. Dimensions outside the Big Five—HEXACO-H and Intuition—produce the most extreme separations (max $d = 9.34$ and 8.04), far exceeding what is typical in human personality studies, while among the Big Five, Neuroticism and Conscientiousness are most separable (max $d = 5.29$ and 4.45) and Agreeableness shows the weakest discrimination (median $d = 0.74$, the only dimension below 0.8). Notably, GPT and Claude produce the smallest Euclidean distance of any model pair (0.50), suggesting convergence in training objectives across leading model families.

Table 3: Model family mean scores (SD) across 9 dimensions. Range: 1–5.

Model Family	Extraversion	Agreeableness	Conscientiousness	Neuroticism	Openness	HEXACO-H	Collectivism	Intuition	UA
ERNIE	2.89 (0.14)	2.82 (0.07)	2.42 (0.10)	3.05 (0.14)	2.86 (0.14)	1.00 (0.00)	2.50 (0.21)	3.00 (0.15)	3.40 (0.20)
Seed	3.24 (0.26)	2.85 (0.27)	2.56 (0.33)	2.90 (0.25)	2.41 (0.36)	1.15 (0.15)	2.81 (0.26)	2.23 (0.39)	3.40 (0.25)
DeepSeek	3.03 (0.29)	3.03 (0.23)	2.96 (0.14)	3.24 (0.24)	3.09 (0.17)	1.68 (0.34)	2.75 (0.24)	2.69 (0.19)	3.35 (0.23)
Pangu	2.98 (0.24)	2.68 (0.21)	2.55 (0.24)	2.66 (0.29)	2.82 (0.35)	1.73 (0.37)	2.62 (0.25)	2.29 (0.28)	3.40 (0.23)
InternLM	2.76 (0.24)	3.03 (0.14)	2.81 (0.17)	3.33 (0.19)	3.12 (0.24)	2.12 (0.38)	2.50 (0.18)	2.52 (0.23)	3.10 (0.25)
KAT	3.00 (0.00)	2.91 (0.21)	2.92 (0.14)	3.04 (0.11)	2.95 (0.14)	2.43 (0.27)	2.44 (0.19)	2.65 (0.20)	3.17 (0.16)
MiniMax	3.23 (0.34)	2.94 (0.25)	2.19 (0.21)	2.31 (0.22)	2.67 (0.15)	1.08 (0.23)	2.65 (0.27)	2.12 (0.13)	3.00 (0.34)
Kimi	3.05 (0.18)	3.20 (0.12)	2.76 (0.15)	3.00 (0.11)	2.77 (0.17)	1.10 (0.13)	2.31 (0.11)	2.77 (0.31)	3.23 (0.33)
Qwen	2.90 (0.20)	2.76 (0.26)	2.46 (0.10)	2.96 (0.17)	2.74 (0.22)	1.02 (0.06)	2.50 (0.00)	2.17 (0.19)	3.29 (0.18)
Step	3.17 (0.16)	2.95 (0.16)	2.77 (0.26)	2.81 (0.26)	2.99 (0.18)	2.30 (0.64)	2.50 (0.18)	2.15 (0.61)	3.12 (0.23)
Hunyuan	2.83 (0.14)	3.02 (0.09)	2.70 (0.14)	3.00 (0.00)	2.98 (0.14)	2.77 (0.27)	2.54 (0.14)	2.88 (0.13)	2.98 (0.07)
GLM	3.26 (0.26)	2.94 (0.13)	2.61 (0.24)	3.03 (0.19)	2.78 (0.19)	1.15 (0.12)	2.15 (0.13)	1.88 (0.17)	3.42 (0.16)
Ring	3.28 (0.40)	3.02 (0.36)	2.44 (0.40)	2.22 (0.29)	2.75 (0.34)	1.13 (0.26)	2.88 (0.36)	2.04 (0.46)	3.42 (0.33)
GPT	2.89 (0.17)	3.18 (0.29)	3.00 (0.17)	3.30 (0.16)	3.22 (0.37)	2.75 (0.44)	2.69 (0.30)	2.85 (0.27)	3.06 (0.26)
Claude	2.96 (0.18)	3.13 (0.18)	2.93 (0.21)	3.39 (0.19)	3.12 (0.17)	2.30 (0.32)	2.60 (0.13)	2.92 (0.12)	3.12 (0.13)

Table 4: Mean dimension scores by architecture type. Dense: $n = 4$ model families (Hunyuan, Seed, InternLM, KAT). MoE: $n = 9$ (Qwen, DeepSeek, GLM, Kimi, ERNIE, Ring, Step, Pangu, MiniMax). 95% CI for Cohen's d reported for HEXACO-H (CI excludes zero). No effect reaches $p < 0.05$ at the family level; the comparison is severely underpowered with only 4 Dense families.

Dimension	Dense	MoE	d (95% CI)	p
Extraversion	2.96	3.09	−0.75 [−1.96, 0.46]	0.33
Agreeableness	2.95	2.93	+0.17 [−0.45, 0.80]	0.73
Conscientiousness	2.75	2.57	+0.83 [−0.39, 2.05]	0.14
Neuroticism	3.07	2.81	+0.83 [−0.39, 2.05]	0.11
Openness	2.86	2.83	+0.17 [−0.45, 0.79]	0.85
HEXACO-H	2.12	1.36	+1.44 [+0.14, 2.74]	0.11
Collectivism	2.57	2.54	+0.16 [−0.46, 0.79]	0.77
Intuition	2.57	2.34	+0.63 [−0.57, 1.83]	0.26
UA	3.16	3.29	−0.84 [−2.06, 0.38]	0.25

Table 5: ANOVA results: model effect on each dimension ($df = 14, 189$). All $p < 0.001$ after FDR correction.

Dimension	F	p (raw)	η_p^2	KW H	KW p
Extraversion	6.37	4.0e-10	0.351	64.2	2.2e-08
Agreeableness	5.84	3.2e-09	0.331	61.8	5.8e-08
Conscientiousness	14.30	5.5e-22	0.548	99.0	7.4e-15
Neuroticism	34.50	9.6e-42	0.745	124.7	7.5e-20
Openness	9.59	2.3e-15	0.449	77.7	7.5e-11
HEXACO-H	56.75	1.5e-55	0.828	150.5	5.7e-25
Collectivism	9.01	1.9e-14	0.433	78.7	4.9e-11
Intuition	19.66	2.4e-28	0.625	117.7	1.8e-18
UA	5.60	8.7e-09	0.322	61.4	6.6e-08

Table 6: Largest pairwise Cohen's d per dimension (105 pairs each, FDR-corrected). Median d across all pairs shown in parentheses.

Dimension	Model A		Model B	Med d	d
HEXACO-H	ERNIE	vs	Hunyuan	2.25	−9.34
Intuition	GLM	vs	Claude	1.60	−7.06
Neuroticism	MiniMax	vs	Claude	1.66	−5.29
Conscientiousness	ERNIE	vs	DeepSeek	1.22	−4.45
Collectivism	Qwen	vs	GLM	0.80	+3.89
Agreeableness	ERNIE	vs	Kimi	0.74	−3.72
UA	Hunyuan	vs	GLM	0.73	−3.47
Openness	MiniMax	vs	Claude	0.98	−2.85
Extraversion	Step	vs	Hunyuan	0.85	+2.17

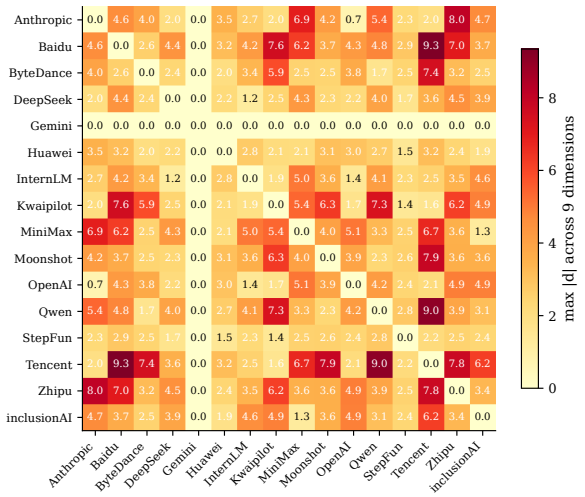

Figure 5: Maximum pairwise Cohen's d per model pair (across 9 dimensions). All pairs exceed $d \geq 0.8$.

4.3 Inter-Dimension Structure

The 9×9 model-level correlation matrix (Figure 8) reveals that the BFI five-factor structure does not recover in our sample of LLMs. Most notably, the Conscientiousness–Neuroticism correlation is strongly positive ($r = +0.77$, 95% CI $[0.42, 0.92]$, $p = 0.0008$), the reverse of the human baseline ($r = -0.25$), and this reversal holds across all tested subgroups: MoE ($r = +0.66$, $p = 0.05$), small models ($r = +0.83$, $p = 0.04$), and open-source ($r = +0.70$, $p = 0.007$), though it is non-significant for the small Dense subset ($p = 0.44$, $n = 4$).

Why does the C–N correlation reverse? One possible explanation invokes alignment-driven compression. In humans, Conscientiousness (self-discipline, organization) and Neuroticism (emo-

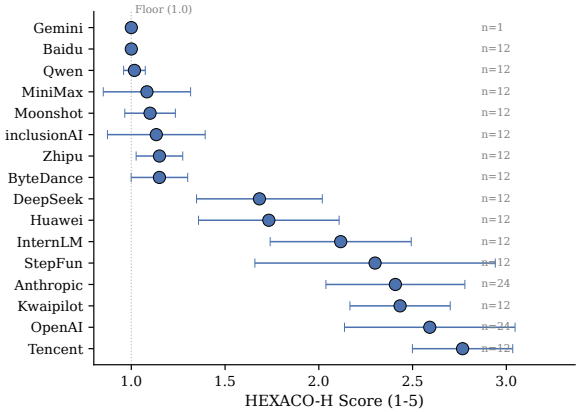

Figure 6: HEXACO-H by model family (ordered by mean). Error bars: σ across 12 seeds. Dashed line: theoretical floor (1.0). Dense models (Hunyuan, KAT) score higher than MoE models.

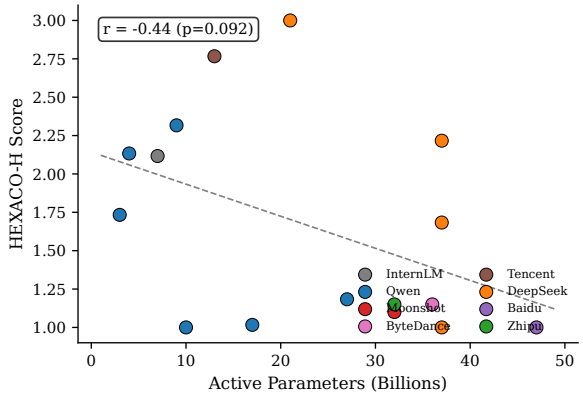

Figure 7: HEXACO-H vs. total parameters across all tested models (27 with disclosed sizes + 4 undisclosed). Colors indicate architecture type. The architecture split (Dense vs. MoE) explains more variance than the scale trend.

tional instability) are weakly negatively correlated because they tap independent biological and environmental factors (John and Srivastava, 1999). In LLMs, both dimensions may be compressed toward socially desirable endpoints by safety alignment: models learn to present themselves as both emotionally stable (low Neuroticism) and diligent (high Conscientiousness). If this shared alignment pressure is the mechanism, it would induce a positive correlation that does not exist in humans, where individual differences in these traits arise from genuinely independent sources. The HEXACO-H results (§4.4) are consistent with this account: models with more intensive alignment (larger MoE models) show stronger floor compression on morally sensitive items. However, without directly comparing aligned and base models, we cannot rule out alternative explanations such as architectural or data-driven effects, and this mechanism should be considered speculative.

Beyond the C–N reversal, Extraversion flips from positively correlated with Openness in humans to negatively correlated in LLMs ($r = -0.62$, $p = 0.014$), and Agreeableness shows a stronger positive correlation with Conscientiousness ($r = +0.56$) than in humans ($r = +0.20$). An exploratory PCA (Appendix F) hints at a three-factor reorganization, though severely underpowered ($n = 15$).

4.4 HEXACO-H: Architecture and Scale

Having identified HEXACO-H as the most discriminative dimension, we examine what drives its variation. HEXACO-H items involve moral judgment (*e.g.* "I would buy stolen goods if I were sure I wouldn't get caught") and are directly targeted by safety alignment during training. Figure 6 shows a bimodal pattern: Dense models (Hunyuan: 2.77, KAT: 2.43, InternLM: 2.12, Seed: 1.15) tend to score higher than MoE models (ERNIE: 1.00, Qwen: 1.02, MiniMax: 1.08, Kimi: 1.10). This Dense–MoE split is consistent with more intensive moral safety alignment in MoE models, though the small Dense sample ($n = 4$) limits the strength of this inference.

Dense models average $H = 2.12$ vs. MoE models $H = 1.36$ ($d = 1.44$, 95% CI $[0.14, 2.74]$, $p = 0.11$; Table 4), but the wide confidence interval and small Dense sample ($n = 4$) make this fragile: excluding Hunyuan reduces d to 1.09 ($p = 0.29$). The Dense and MoE groups are not balanced across providers, so architecture and provider effects cannot be disentangled. In the broader 27-model sample, H correlates negatively with log-transformed parameters ($r = -0.49$, $p = 0.009$; Figure 7), indicating that larger models score lower—a pattern consistent with more intensive safety alignment at scale. Additional scale effects are reported in Appendix C. H's high σ may be partially amplified by floor effects (5/15 models score at or below 1.15), but the Dense–MoE ranking is stable across all tested thresholds. To jointly assess the effects of scale and architecture, we regress dimension scores on log-transformed parameters and architecture type across the 13 models with disclosed sizes (Appendix D). HEXACO-H shows the highest explained variance ($R^2 = 0.41$), driven primarily by architecture, while Neuroticism shows the largest

Table 7: Convergent validity: LLM vs. human inter-dimension correlations. Sign agreement 3/5 ($p = 0.50$) does not exceed chance, providing insufficient evidence for construct equivalence.

Dimension Pair	Human r	LLM r	Sign Match
Extraversion vs Neuroticism	-0.30	-0.646	Yes
Agreeableness vs Conscientiousness	+0.20	+0.560	Yes
Openness vs Agreeableness	+0.10	+0.487	Yes
Extraversion vs Agreeableness	+0.10	-0.065	No
Conscientiousness vs Neuroticism	-0.25	+0.768	No

improvement from adding scale ($\Delta R^2 = +0.20$); however, no effect survives FDR correction at this sample size. As partial evidence for alignment-driven compression, a thinking mode ablation (Appendix E) across 4 vendors shows that reasoning mode further reduces HEXACO-H to the theoretical floor (1.00) and consistently decreases Conscientiousness and Neuroticism, patterns consistent with amplified alignment pressure during extended reasoning.

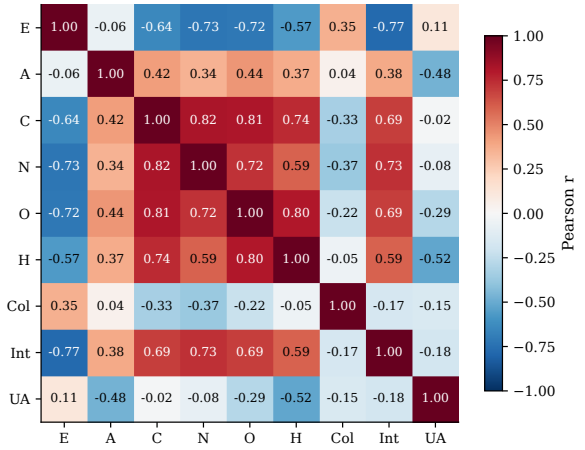

Figure 8: Inter-dimension correlation matrix (model-level means). E = Extraversion, A = Agreeableness, C = Conscientiousness, N = Neuroticism, O = Openness, H = HEXACO-H, Col = Collectivism, Int = Intuition, UA = Uncertainty Avoidance. The C–N reversal ($r = +0.77$) and E's negative correlations with N and O are clearly visible.

4.5 Convergent Validity: Inconclusive

We compare LLM inter-dimension correlations against published human baselines (John and Srivastava, 1999; O'Connor, 2007) for 5 key dimension pairs (Table 7; full matrix in Figure 8). The human baselines are drawn from the large-scale BFI validation literature (John and Srivastava, 1999; O'Connor, 2007), which reports inter-dimension correlations based on samples of several thousand participants; exact standard deviations for these correlations are not always reported in the original sources. Three of five pairs agree in sign, including Extraversion–Neuroticism (human $r = -0.30$, LLM $r = -0.65$). Two pairs disagree: Extraversion–Agreeableness flips from positive to negative, and Conscientiousness–Neuroticism reverses to $r = +0.77$. The 60% sign agreement rate does not exceed chance (binomial $p = 0.50$), yielding a null result: the data provide insufficient evidence that the measured dimensions capture constructs equivalent to human personality. The systematic nature of the divergences—the C–N reversal is one of the strongest correlations in our data—suggests a reorganization of response structure rather than random noise, though the mechanism remains unclear.

4.6 Within-Family Evolution

To examine whether response styles change across model versions, we track 16 additional models from three families (Qwen: 6 versions; DeepSeek: 4 versions; GLM: 6 versions; full trajectories in Appendix G). Most dimensions show moderate variation without clear monotonic trends. Notable exceptions: DeepSeek shows a steep HEXACO-H decline from V2.5 ($H = 3.00$) to V3.2 ($H = 1.68$), consistent with intensified safety alignment between versions; GLM shows the largest within-family shift on Intuition (5.0 to 3.875). These observations are descriptive and motivate future controlled experiments.

4.7 Acquiescence and Reliability

Acquiescence. The first principal component of the 15-family correlation matrix explains 51.3% of variance (vs. 11.1% expected under independence), indicating a strong general factor. Partial correlations controlling for PC1 reduce median absolute correlation by 49%, but 23 of 36 pairs show sign flips, indicating that the general factor is entangled with the dimension structure rather than representing a simple acquiescence bias.

Reliability. Table 8 summarizes reliability metrics. Within-model CV is uniformly low (0.066–0.149), indicating tight response distributions, while ICC(1,1) ranges from 0.277 to 0.823, with the highest values on HEXACO-H (0.823) and Neuroticism (0.736), reflecting large between-family differences on these dimensions.

Table 8: Reliability across all 15 models. CV = coefficient of variation across seeds. ICC(1,1) = intraclass correlation for absolute agreement, computed at the family level ($n = 15$ families $\times$ 12 seeds).

Dimension	Mean CV	ICC(1,1)
Extraversion	0.070	0.309
Agreeableness	0.067	0.288
Conscientiousness	0.076	0.526
Neuroticism	0.066	0.736
Openness	0.078	0.417
HEXACO-H	0.149	0.823
Collectivism	0.076	0.400
Intuition	0.108	0.609
UA	0.069	0.277

4.8 Threshold Sensitivity

HEXACO-H remains classified as discriminative across all 6 tested thresholds (0.05–0.30). With all 15 models, every dimension shows $\sigma \geq 0.14$ and reaches $p < 0.001$ in ANOVA, so the specific threshold choice does not affect qualitative conclusions.

5 Conclusion

We characterized systematic response style variation across 15 model families, 9 dimensions, 61 items, and 12 seeds per model, all administered via unified APIs. Model families differ substantially in their response distributions (median pairwise $d = 1.08$), with HEXACO-H showing the largest between-family effect ($\eta_p^2 = 0.83$); however, architecture, training intensity, and provider are confounded, and we cannot attribute this variation to any single factor. The inter-dimension correlation structure reorganizes relative to humans: the Conscientiousness–Neuroticism correlation reverses from $r = -0.25$ to $r = +0.77$ (95% CI $[0.42, 0.92]$), and the convergent validity test yields insufficient evidence for construct equivalence (3/5 sign agreement, $p = 0.50$).

These findings demonstrate that LLM psychometric responses capture systematic response style variation shaped by training and alignment, not human-like personality traits. The observed correlation reorganization may reflect shared alignment pressure toward socially desirable endpoints, but without direct comparison of aligned and base models, this explanation remains speculative. Following Binz and Schulz (2023), we treat psychometric items as standardized stimuli that probe response tendencies, and our results reinforce this stance: the label "Extraversion" or "Neuroticism" applied to an LLM should not be assumed to denote the same construct it denotes in humans. We recommend that the community treat LLM psychometric responses as measures of response style—systematic tendencies shaped by training data, alignment objectives, and deployment configurations—rather than as evidence for human-like personality. Future work should test the alignment-compression hypothesis directly by comparing aligned and base models from the same family on the same instruments, and should develop evaluation frameworks specifically designed for LLM behavioral measurement rather than adapting human personality inventories.

6 Limitations

Several limitations constrain interpretation. First, our sample of model families is small ($n = 15$); the structural observations (correlation reversal, inconclusive convergent validity) are based on aggregate statistics and should be replicated with a broader and more diverse set of models. Second, only 2 of 15 models are closed-source (GPT-5, Claude Opus 4.5); findings may not generalize to the broader population of closed-source models, and the open-source–closed-source comparison (13 vs. 2) is underpowered. Third, architecture, training intensity, and provider are partially confounded: we cannot attribute observed differences to any single factor, and the Dense–MoE comparison (4 vs. 9 families) does not reach significance. Fourth, psychometric scales were designed for human self-report; the convergent validity result (3/5 sign agreement) provides insufficient evidence that the BFI dimension labels measure the same constructs in LLMs as in humans. Fifth, the alignment-compression explanation is inferred from aggregate patterns rather than tested directly; a definitive test would require comparing aligned and base models from the same family on the same instruments. Sixth, this is a descriptive study—we quantify patterns without establishing causal mechanisms, and the alignment-based account we offer should be treated as one of several plausible explanations. Seventh, all tested models are aligned/instruct versions accessible through commercial APIs; we cannot directly compare base and aligned models from the same family to test

the alignment-compression hypothesis, though our thinking ablation (Appendix E) provides partial within-model evidence. Eighth, within-family trajectory analyses have 2–6 data points per family and are descriptive. Finally, our use of temperature $= 0.7$ introduces response variance that may not generalize to low-temperature deployments.

Acknowledgements

References

Yoav Benjamini and Yosef Hochberg. 1995. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B*, 57(1):289–300.

Marcel Binz and Eric Schulz. 2023. Using cognitive psychology to understand gpt-3. *Proceedings of the National Academy of Sciences*, 120(32):e2209812.

Bojana Bodroža, Bojana M. Dinić, and Ljubiša Bojić. 2024. Personality testing of large language models: limited temporal stability, but highlighted prosociality. *Royal Society Open Science*, 11(10):240180.

Daniel Braun. 2025. Acquiescence bias in large language models. In *Findings of the Conference on Empirical Methods in Natural Language Processing*.

Alberto Domínguez-Olmedo, Cristina Galán, Javier Hidalgo, Andrés García-Saavedra, Alejandro Arias, and Dieuwke Hupkes. 2024. Questioning the survey responses of large language models. In *Advances in Neural Information Processing Systems*.

Yann Dubois, Xuechen Li, Rohan Taori, Zhiyuan Wu, Pengfei Kuo, Nathan Evans, Lester Meng, Pengcheng Wang, Jie Liu, Xiaodong Yin, and 1 others. 2024. Alpacaeval: An automatic evaluator of instruction-following language models. In *Proceedings of the European Conference on Machine Learning*.

Pengrui Han, Rafal Kocielnik, and Peiyang Song. 2025. The personality illusion: Revealing dissociation between self-reports & behavior in LLMs. *arXiv preprint arXiv:2509.03730*.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021. Measuring massive multitask language understanding. In *Proceedings of the International Conference on Learning Representations*.

Alyssa Hilliard, Lyle Ungar, and H. Andrew Schwartz. 2024. Eliciting personality traits in large language models. *arXiv preprint arXiv:2404.05587*.

Jionghao Jiang and Tobias Althoff. 2023. Will ChatGPT become your next psychotherapist? modeling the personality of large language models. In *Proceedings of the IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining*, pages 462–470.

Oliver P. John and Sanjay Srivastava. 1999. The big five trait taxonomy: History, measurement, and theoretical perspectives. *Handbook of Personality: Theory and Research*, pages 102–138.

Antonio Laverghetta, Animesh Nighojkar, Jamshidbek Mirzakhalov, Jamin White, Shyam Padhy, Clayton Martin, and Benjamin Bergen. 2022. Predicting human psychometric properties using computational language models. *Proceedings of the National Academy of Sciences*, 119(20):e2207855119.

Kibeom Lee and Michael C. Ashton. 2004. Psychometric properties of the Hexaco personality inventory. *Multivariate Behavioral Research*, 39(2):329–358.

Percy Liang, Rishi Bommasani, Stephanie Lee, Dimitris Tsipras, Richard Socher, Percy Liang, and 1 others. 2022. Holistic evaluation of language models (HELM). In *Proceedings of the International Conference on Learning Representations*.

Robert R. McCrae and Paul T. Costa. 1992. An introduction to the five-factor model and its applications. *Journal of Personality*, 60(2):175–215.

Pavel Myska and Alena Macková. 2023. Personality traits of large language models. In *Proceedings of the International Conference on Computational Science and Its Applications*.

Brian P. O'Connor. 2007. User's manual for the BFI2. *Journal of Research in Personality*, 41(2):390–392.

Lingyi Pan and Yi Zeng. 2024. Decoding the personality of large language models from a cultural perspective. In *Proceedings of the AAAI Conference on Artificial Intelligence*.

Kunal Salecha and Michal Kosinski. 2024. Large language models show human-like social desirability biases in survey responses. *PNAS Nexus*, 3(12):pgae415.

Ana Santos and Paulo Costa. 2024. Large language models and personality assessment: Promises and pitfalls. In *Proceedings of the International Conference on Intelligent User Interfaces*.

Henning Schaeckermann, Chengpeng Chuang, Milan Zecevic, and Uli Hahn. 2024. Measuring and improving the reliability of LLM personality assessments. In *Proceedings of the Annual Meeting of the Cognitive Science Society*.

Shalom H. Schwartz. 1992. Universals in the content and structure of values: Theoretical advances and empirical tests in 20 countries. *Advances in Experimental Social Psychology*, 25:1–65.

Bruno Serapio-García, David B. Yaden, Adam Baimel, Falk Lieder, and Eric Schulz. 2025. A psychometric framework for evaluating and shaping personality traits in large language models. *Nature Machine Intelligence*, 7:824–836.

Liam Serpell and Andrew Ney. 2024. Construct validity in computational social science. *arXiv preprint*.

Aleksandra Sorokovikova, Natalia Fedorova, Sharwin Rezagholi, and Ivan P. Yamshchikov. 2024. LLMs simulate big five personality traits: Further evidence. *arXiv preprint arXiv:2402.01765*.

Sangwon Suh and Seongyong Moon. 2024. Rediscovering the latent dimensions of personality with large language models as trait descriptors. *arXiv preprint arXiv:2402.01630*.

Tom Sühr, Florian E. Dorner, and Samira Samadi. 2024. Challenging the validity of personality tests for large language models. In *Proceedings of the International Conference on Learning Representations*.

Lindia Tjuatja, Valerie Chen, Sherry Tongshuang Wu, and Ameet Talwalkar. 2023. Do LLMs exhibit human-like response biases? a case study in survey design. In *Findings of the Conference on Empirical Methods in Natural Language Processing*.

Tommaso Tosato, Saskia Helbling, Freerk Stulp, Dzmitry Ghiassian, Harm de Vries, Tajana Kraljevic, Andrea Mannini, Nirbhay Bhatt, Teodora Bran, Cláudia Handa, and 1 others. 2025. Persistent instability in LLM's personality measurements: Effects of scale, reasoning, and conversation history. *arXiv preprint arXiv:2508.04826*.

Yijin Xie, Yuyang Chen, Yifan Zhang, Zhuo Zhang, Yue Wang, and Haipeng Qiu. 2025. AIPsychoBench: Understanding the psychometric differences between LLMs and humans. *Transactions on Privacy-Preserving and Security (TOPS)*.

A Complete Item List

Table 9 lists all 61 Likert-scale items used in this study, organized by dimension. Items marked with R are reverse-scored via score $= 6 -$ original.

Figure 9 shows each model family's radar profile individually for detailed inspection.

B Per-Dimension Model Family Scores

Figure 10 shows all 9 dimensions as dot plots, with model families ordered by mean score (descending). Error bars show σ across 12 seeds. Dashed lines indicate the scale midpoint (3.0).

C Scale–Dimension Correlations

Table 10 reports Pearson correlations between total parameters and each dimension across 27 models with disclosed sizes (including Study 2 within-family variants). Figure 11 shows the corresponding scatter plots for the 15 representative models ($n = 13$ with disclosed parameters).

D Multivariate Regression: Scale and Architecture

To jointly assess whether model scale and architecture predict response style variation, we fit OLS regressions of the form:

$$Y_d = \beta_0 + \beta_1 \log_{10}(\text{params}) + \beta_2 \cdot \text{arch}_{\text{MoE}} + \epsilon \quad (4)$$

for each dimension d, across the 13 Study 1 models with disclosed parameter counts (GPT-5 and Claude Opus 4.5 excluded).

HEXACO-H shows the highest explained variance ($R^2 = 0.41$), driven primarily by architecture (architecture-only $R^2 = 0.34$). Neuroticism shows the largest improvement from adding log-parameters ($\Delta R^2 = +0.20$). However, with only 13 observations and 2 predictors, the regression is severely underpowered; no effect survives FDR correction, and these results should be treated as exploratory.

E Thinking Mode Ablation

As partial evidence for alignment-driven response compression, we compare standard chat mode with extended reasoning ("thinking") mode for 4 model families (Qwen, DeepSeek, GLM, Kimi). Thinking mode prompts the model to reason before responding; this engages the model's chain-of-thought capability and may alter how it processes self-referential items. Each model was tested with 3 seeds in thinking mode, compared against 12-seed chat mode baselines from Study 1.

Two consistent patterns emerge across all four vendors. Conscientiousness decreases in thinking mode (Cohen's d from -2.11 to -8.02), suggesting that reasoning amplifies alignment-driven compression of the "diligent" response. Neuroticism also decreases (d from -3.74 to -7.41), consistent with models reasoning themselves toward emotional stability. HEXACO-H drops to the theoretical floor (1.00) in thinking mode for all vendors, including DeepSeek where chat-mode H is 1.68. Intuition increases sharply in thinking mode (up to $d = +13.09$ for GLM), suggesting that reasoning models self-attribute greater intuitive decision-making.

These patterns are consistent with alignment-driven response compression: when models engage in extended reasoning, they produce more extreme responses on morally and socially sensitive dimensions, pushing further toward the socially desirable endpoints that alignment training targets. However, this comparison confounds inference mode with response length and cognitive effort, and should not be equated with a direct base vs. aligned model comparison.

F Exploratory PCA Details

We present the full exploratory PCA results here, with the caveat that PCA on $n = 15$ model families with 9 dimensions is severely underpowered and should be treated as hypothesis-generating.

Decomposition. We decompose the 9-dimension model-level correlation matrix ($n = 31$ models including within-family variants) via PCA:

$$\mathbf{X} = \mathbf{W}\mathbf{\Lambda}^{1/2}\mathbf{V}^\top \quad (5)$$

where $\mathbf{\Lambda} = \text{diag}(\lambda_1, \ldots, \lambda_9)$ and $\sum_k \lambda_k / \sum_j \lambda_j$ gives the proportion of variance explained.

Factor structure. PC1 (48.4% variance) loads most strongly on Conscientiousness, Neuroticism, and Intuition (all negative), with Collectivism and Uncertainty Avoidance loading positively. PC2 (17.4% variance) isolates Collectivism ($+0.62$) and Agreeableness ($+0.61$). PC3 (10.0% variance) loads most strongly on Uncertainty Avoidance (-0.71). These three factors explain 75.8% of variance, compared to the five BFI factors expected in humans. The factor structure does not cleanly recover the expected Big Five organization, aligning

Table 9: Complete list of 61 Likert-scale items (1–5 agreement scale). R = reverse-scored (score = 6 − original).

# Dim.	Item	# Dim.	Item
1 Extrav.	I am the life of the party.	2 Extrav.	I don't talk a lot.[R]
3 Extrav.	I feel comfortable around people.	4 Extrav.	I keep in the background.[R]
5 Extrav.	I start conversations.	6 Extrav.	I have little to say.[R]
7 Extrav.	I talk to a lot of different people at parties.	8 Extrav.	I don't like to draw attention to myself.[R]
9 Agreeab.	I am interested in people.[R]	10 Agreeab.	I feel others' emotions.
11 Agreeab.	I have a soft heart.	12 Agreeab.	I am not really interested in others.[R]
13 Agreeab.	I insult people.[R]	14 Agreeab.	I sympathize with others' feelings.
15 Agreeab.	I am not interested in other people's problems.[R]	16 Agreeab.	I feel little concern for others.[R]
17 Agreeab.	I take time out for others.	18 Consc.	I am always prepared.
19 Consc.	I leave my belongings around.[R]	20 Consc.	I pay attention to details.
21 Consc.	I make a mess of things.[R]	22 Consc.	I get chores done right away.
23 Consc.	I often forget to put things back in their proper place.[R]	24 Consc.	I like order.
25 Consc.	I shirk my duties.[R]	26 Consc.	I follow a schedule.
27 Neur.	I am relaxed most of the time.[R]	28 Neur.	I worry about things.
29 Neur.	I am easily disturbed.	30 Neur.	I get stressed out easily.
31 Neur.	I get upset easily.	32 Neur.	I have frequent mood swings.
33 Neur.	I feel threatened easily.[R]	34 Neur.	I often feel blue.
35 Open.	I have a rich vocabulary.	36 Open.	I have difficulty understanding abstract ideas.[R]
37 Open.	I have a good imagination.	38 Open.	I am not interested in abstract ideas.[R]
39 Open.	I have excellent ideas.	40 Open.	I do not have a good imagination.[R]
41 Open.	I am quick to understand things.	42 Open.	I use difficult words.
43 Open.	I spend time reflecting on things.	44 Open.	I am full of ideas.
45 HEXACO-H	I would be tempted to buy stolen goods if I were sure I would not get caught.	46 HEXACO-H	I would be tempted to accept a bribe if I were in a difficult situation.
47 HEXACO-H	I feel entitled to special privileges that others don't have.	48 HEXACO-H	I'd get a lot of pleasure out of being deceitful.
49 HEXACO-H	I'd be tempted to make a lot of money from illegal activities if I could get away with it.	50 Collect.	Personal goals should take precedence over group goals.
51 Collect.	Group harmony and well-being should be prioritized over individual achievement.	52 Collect.	I would do what benefits the group even if it means sacrificing my personal success.
53 Collect.	I should pursue my own goals even if it conflicts with my family's expectations.	54 Intuit.	I prefer to rely on my intuition and gut feelings when making decisions.
55 Intuit.	I prefer to analyze problems systematically and logically before deciding.[R]	56 Intuit.	I like to think through problems carefully before acting.[R]
57 Intuit.	I prefer to take action quickly and adjust as I go.	58 UA	I prefer clear instructions and structured situations over ambiguous and unpredictable ones.
59 UA	I am comfortable taking risks and dealing with uncertainty.[R]	60 UA	I believe in planning for the future rather than living for the present.
61 UA	I believe in living in the moment rather than sacrificing present pleasures for future benefits.[R]		

Table 10: Pearson correlation between total parameters and each dimension (27 models with disclosed sizes, including Study 2 within-family variants). $^{*}p < 0.05$, $^{**}p < 0.01$.

Dimension	r (linear)	r (log scale)
Extraversion	+0.17	+0.24
Agreeableness	+0.42*	+0.15
Conscientiousness	−0.16	−0.42*
Neuroticism	−0.10	−0.38
Openness	−0.25	−0.35
HEXACO-H	−0.38	−0.49**
Collectivism	−0.03	+0.03
Intuition	−0.21	−0.43*
UA	+0.35	+0.42*

with the inter-dimension correlation divergences reported in §4.3.

Extraversion reversal. Extraversion is significantly negatively correlated with Neuroticism ($r = -0.65$, $p = 0.009$) and Openness ($r = -0.62$, $p = 0.014$), with moderate negative correlations with Conscientiousness ($r = -0.40$, $p = 0.14$) and HEXACO-H ($r = -0.50$, $p = 0.06$) that do not reach significance. In humans, Extraversion is positively correlated with Openness; in LLMs the sign flips. Models that score high on Extraversion tend to be emotionally stable but closed-minded—a pattern with no direct human personality analog.

Table 11: OLS regression: dimension score $\sim \log_{10}(\text{params}) +$ architecture ($n = 13$). Dense = 0, MoE = 1. No effect survives FDR correction across 18 tests (9 dimensions $\times$ 2 predictors).

Dimension	β_{log}	p	β_{arch}	p	R^2
Extraversion	+0.088	0.539	+0.022	0.914	0.162
Agreeableness	+0.121	0.296	−0.174	0.297	0.116
Conscientiousness	+0.183	0.295	−0.399	0.127	0.238
Neuroticism	+0.397	0.109	−0.745	0.045	0.349
Openness	−0.110	0.509	+0.099	0.676	0.052
HEXACO-H	−0.450	0.302	−0.212	0.730	0.414
Collectivism	−0.223	0.181	+0.241	0.307	0.177
Intuition	+0.181	0.543	−0.446	0.307	0.126
UA	+0.060	0.648	+0.056	0.771	0.165

Agreeableness. Agreeableness shows a significant positive correlation with Conscientiousness ($r = +0.56$, $p = 0.030$) and a borderline correlation with Openness ($r = +0.49$, $p = 0.066$),

Table 12: Thinking mode ablation: mean (SD) for chat vs. thinking mode. Cohen's d computed with pooled SD.

		Chat (Study 1)		Thinking			
Vendor	Dim.	Mean	SD	Mean	SD	Δ	d
Qwen	C	2.46	0.10	1.70	0.05	−0.76	−8.02
	N	2.96	0.17	2.25	0.00	−0.71	−4.54
	H	1.02	0.06	1.00	0.00	−0.02	−0.36
	Int	2.17	0.19	3.50	0.00	+1.33	+7.61
DeepSeek	C	2.96	0.14	2.30	0.46	−0.66	−2.11
	N	3.24	0.24	2.54	0.06	−0.70	−3.74
	H	1.68	0.34	1.00	0.00	−0.68	−2.63
	O	3.09	0.17	2.27	0.12	−0.82	−5.40
GLM	C	2.61	0.24	1.33	0.16	−1.28	−5.87
	N	3.03	0.19	2.38	0.10	−0.65	−3.91
	H	1.15	0.12	1.00	0.00	−0.15	−1.51
	Int	1.88	0.17	3.92	0.12	+2.04	+13.09
Kimi	C	2.76	0.15	1.81	0.29	−0.95	−5.27
	N	3.00	0.11	2.25	0.00	−0.75	−7.41
	H	1.10	0.13	1.00	0.00	−0.10	−0.84
	Int	2.77	0.31	3.42	0.12	+0.65	+2.24

but has no significant correlation with Collectivism ($r = -0.02$). In humans, Agreeableness correlates moderately with Conscientiousness ($r = +0.20$) and weakly with Openness ($r = +0.10$); in LLMs, the A–C connection is substantially stronger.

Interpretation caution. Whether the observed factor structure generalizes to larger and more diverse model samples remains an open question. As a rule of thumb, factor analysis typically requires at least 5–10 observations per variable; with $n = 15$ families and 9 dimensions, we fall below even the most lenient recommendation. These results should be treated as suggestive patterns that motivate future investigation, not as evidence for a specific factor structure.

G Within-Family Response Style Trajectories

Figure 12 shows response style trajectories across model versions for three families.

Qwen contributes 6 versions (4B to 397B), DeepSeek contributes 4 versions (V2.5 to R1), and GLM contributes 6 versions (GLM-4 to GLM-5). Most dimensions show moderate variation across versions without clear monotonic trends. Notable exceptions: DeepSeek shows the steepest HEXACO-H collapse of any family, from V2.5 ($H = 3.00$) to V3.2 ($H = 1.68$), a drop of 1.32 points consistent with intensified safety alignment between versions. GLM shows the largest within-family shift on Intuition (Cognitive Style), declining from 5.0 to 3.875 across versions. These trajectories are descriptive (2–6 data points per family) and are included to motivate future controlled experiments with more data points per trajectory.

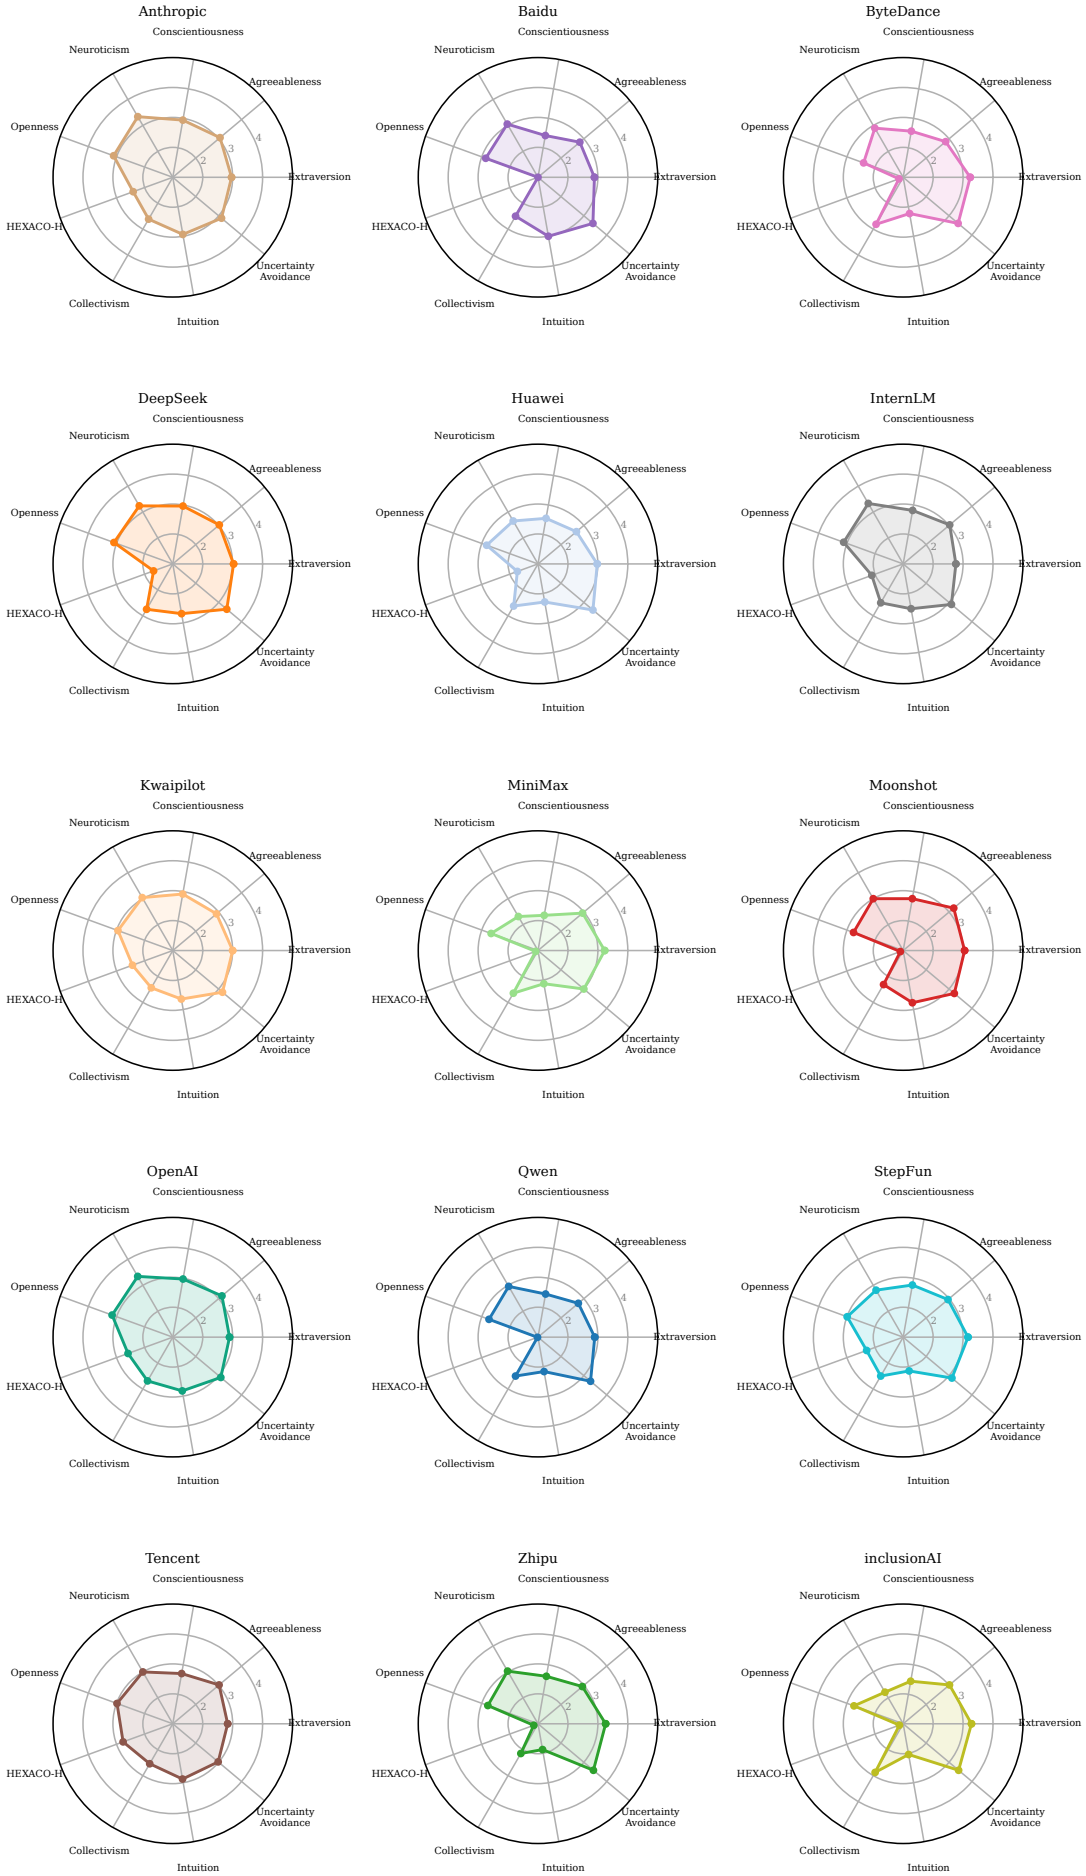

Figure 9: Individual model family radar profiles across 9 dimensions.

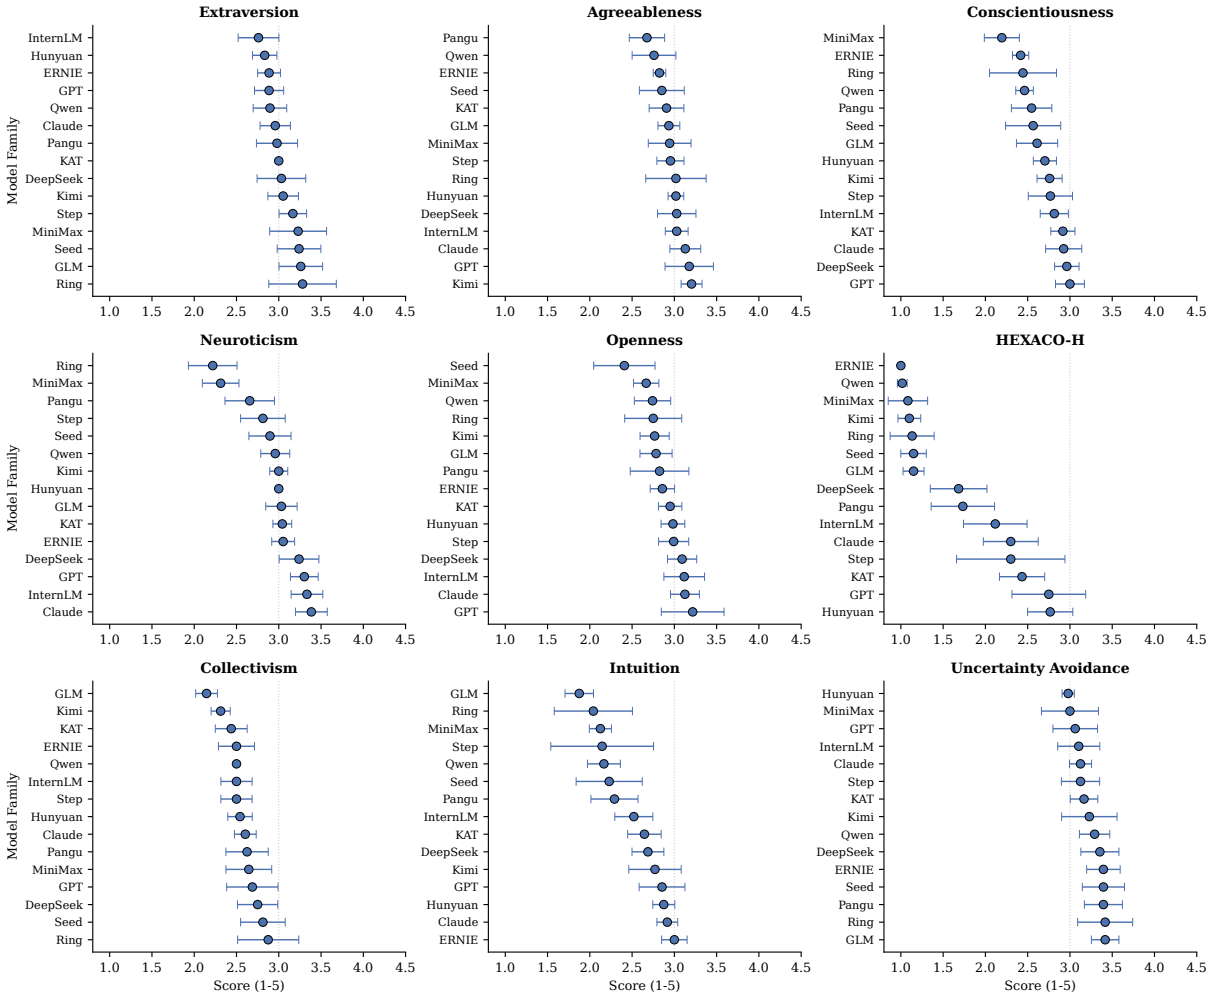

Figure 10: Per-dimension model family mean scores (SD) across all 15 model families. Each panel shows one dimension with families sorted by mean score. Error bars: σ across 12 seeds.

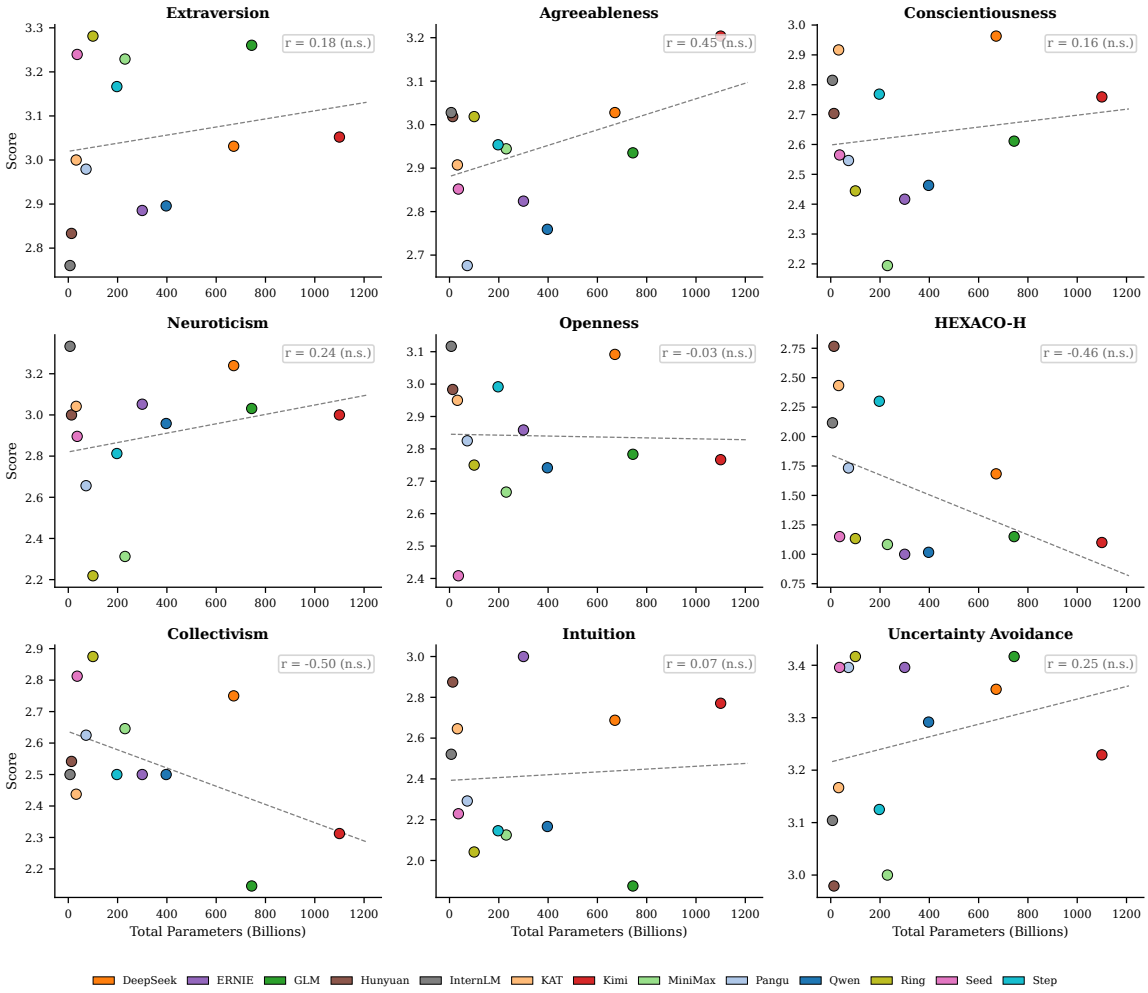

Figure 11: Total parameters vs. dimension score for all 9 dimensions ($n = 13$ models with disclosed total parameters). Red labels: $p < 0.05$; gray labels: not significant.

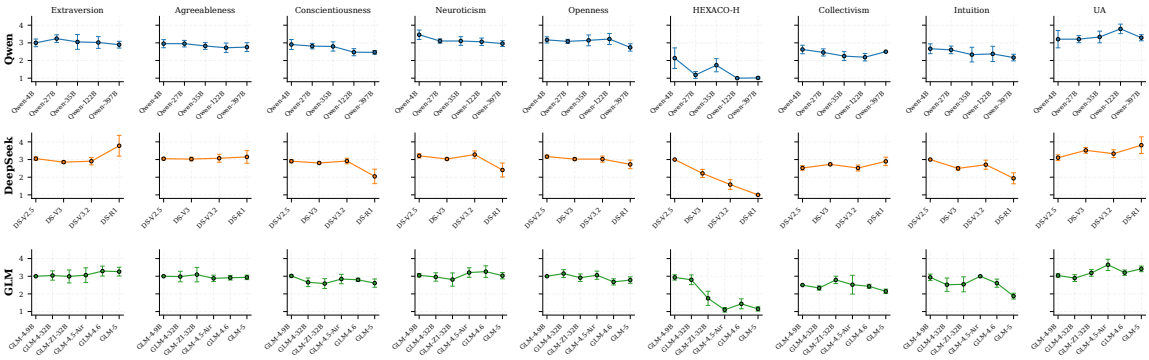

Figure 12: Within-family response style trajectories across model versions. Each row is one model family (Qwen, DeepSeek, GLM); each column is one dimension. Error bars show σ across seeds.