

Acquiescence Is Not Personality: Why Human Psychometric Instruments Fail on LLMs

Anonymous Authors

Abstract

A growing body of work administers human personality questionnaires to LLMs and reports “personality profiles.” We ask a prior question: do these instruments actually measure anything meaningful in LLMs? We gave a 221-item battery (IPIP-NEO-120, SD3, ZKPQ-50-CC, EPQR-A) to 20 LLMs and subjected the 75,140 responses to standard psychometric validation. The core problem is acquiescence: LLMs tend to agree with every statement, regardless of direction. This single bias collapses the measurement. Agreeableness, the domain with the most reverse-coded items, yields $\alpha = 0.06$, because forward and reverse items pull in opposite directions under blanket agreement. The Big Five factor structure collapses from five to three factors. Inter-model variance is below 1%: “personality differences” between models are effectively noise. The instruments are not measuring LLM personality; they are measuring the artifacts of alignment training.

1 Introduction

Can we measure the “personality” of a large language model? A growing body of work thinks so. Researchers administer human personality questionnaires to LLMs, compute scores, and report personality profiles (Miotto et al., 2023; Bhandari et al., 2025; Serapio-García et al., 2025; Sorokovikova et al., 2023; Lee et al., 2025). Some even propose using LLMs as “silicon samples” to simulate survey respondents (Horton, 2023).

But here is the problem. These instruments were designed for humans. A personality questionnaire is only valid for a population if it passes basic checks: internal consistency (Cronbach’s α), correct factor structure, convergent validity, and measurement invariance across subgroups (Cronbach, 1951; Lord and Novick, 1968). Nobody checks these things before interpreting LLM personality scores.




Figure 1: Conceptual framing of the research question. The paper does not assume that questionnaire scores are personality traits; it first asks whether the measurement conditions required for that interpretation hold for LLM responses.

The cracks are already showing. Sühr et al. (2025) found that LLMs endorse both forward-coded and reverse-coded items at the same time. Tjautja et al. (2023) showed that LLM response biases in surveys differ systematically from human patterns. Acerbi and Stubbersfield (2024) and Salecha and Ireland (2024) independently found that LLMs shift responses toward socially desirable endpoints by up to 1.2 human standard deviations. The instability goes deeper: Pelet et al. (2025) showed that reordering questions or changing conversation history produces different personality scores. Huang et al. (2024) found that scale reliability is highly sensitive to instructions and lan-

guage. Han et al. (2025) showed that self-reported personality in LLMs dissociates from actual behavior. Each of these findings chips away at the assumption that the instruments work. But each study looks at a different symptom. Nobody has run all the checks together, on the same data, to find the root cause.

That is what we do. Rather than profiling LLM personality, we ask a simpler question: do these instruments measure anything at all when given to LLMs? We administered a 221-item battery (the IPIP-NEO-120 (Johnson, 2014), the Short Dark Triad (Jones and Paulhus, 2014), the ZKPQ-50-CC (Aluja et al., 2006), and the EPQR-A (Francis et al., 1992)) to 20 LLMs across 17 role-conditioned prompts. Then we ran the same validation checks that psychometricians run before using any instrument on a new population. Every check failed. The cause is acquiescence: LLMs tend to agree with every statement, regardless of direction. This single bias explains why internal consistency collapses in reverse-heavy domains ($\alpha = 0.06$ for Agreeableness), why the Big Five factor structure shrinks to three, and why inter-model variance is below 1%.

What makes our approach different from prior work is that we do not just report that LLM personality measurements are unreliable. We identify the specific mechanism (acquiescence), show which domains it destroys (those with many reverse-coded items), and demonstrate that the resulting scores reflect alignment artifacts rather than genuine traits. We also show a previously unreported asymmetry: normal personality items (IPIP) show positive acquiescence (agree with everything), while antisocial content (SD3) shows negative social desirability (reject dark traits). Two biases, one origin.

2 Related Work

2.1 Personality Measurement in LLMs

Researchers have administered personality questionnaires to LLMs and reported personality profiles (Miotto et al., 2023; Serapio-García et al., 2025; Bhandari et al., 2025; Sorokovikova et al., 2023; Lee et al., 2025; Hagendorff, 2024). But almost none of these studies check whether the instruments actually work on LLMs the way they work on humans. Benchmarks show systematic deviations between LLM and human responses (Xie et al., 2025), temporal stability is limited (Bodroža et al., 2024), and measurement invariance is as-

sumed rather than tested (Sühr et al., 2025). Chen et al. (2025) noted that most existing frameworks skip psychometric validation entirely.

Our contribution. Previous work asks “what personality does this LLM have?” and reports scores. We ask a different question: do these instruments measure anything meaningful at all when given to LLMs? Instead of profiling, we run the full set of standard psychometric validation checks—internal consistency, factor structure, convergent validity, reverse-item consistency, measurement invariance, and variance decomposition—on a single dataset covering 20 models and 221 items.

2.2 Response Style and Acquiescence Bias

Survey respondents do not always answer based on content alone. Acquiescence bias—agreeing with statements regardless of direction—is one of the most studied response styles (Paulhus, 1991; Lord and Novick, 1968; Baumgartner and Steenkamp, 2007). It is especially damaging for instruments that rely on reverse-coded items. Acerbi and Stubbersfield (2024) and Salecha and Ireland (2024) showed that LLMs also exhibit social desirability biases in Big Five surveys, and Responsible NLP Group (2025) confirmed that acquiescence in LLMs is robust across languages.

Our contribution. These studies documented that acquiescence and social desirability exist in LLM responses. Our approach is different: we quantify the measurement damage they cause by decomposing the bias at item level, separating forward-coded and reverse-coded agree rates, and comparing across instruments with different content (normal personality vs. dark traits).

2.3 Alignment Effects on LLM Behavior

RLHF and safety training make LLMs more helpful and harmless, but also more homogeneous (Ouyang et al., 2022; Bai et al., 2022). Sharma et al. (2024) showed that aligned models tend to echo user opinions even when wrong, and Zheng et al. (2025) linked Agreeableness scores to sycophancy in role-playing LLMs. Wang et al. (2025) argued that RLHF encodes a specific notion of acceptability that may override genuine variation.

Our contribution. These studies examined alignment effects on individual behaviors. We take a different angle: we test whether alignment produces a universal response style across model families via

Table 1: Psychometric instruments in the 221-item battery.

Instrument	Items	Format	Domains	Rev.
IPIP-NEO-120	120	Likert-5	5	41
SD3	27	Likert-5	3	5
ZKPQ-50-CC	50	T/F	5	12
EPQR-A	24	Y/N	4	5
Total	221	3 formats	17	63

DIF analysis, and we quantify how far persona injection can push a model from its default behavior (Persona Separation Degree).

3 Methodology

3.1 Psychometric Instruments

We administered a 221-item battery comprising four validated instruments (Table 1). The **IPIP-NEO-120** (Johnson, 2014) contains 120 Likert-5 items measuring the Big Five domains (John and Srivastava, 1999; McCrae and Costa, 1992) through 6 facets each, with 41 reverse-coded items. The **SD3** (Jones and Paulhus, 2014) contains 27 Likert-5 items measuring Machiavellianism, Narcissism, and Psychopathy, with 5 reverse-coded items. The **ZKPQ-50-CC** (Aluja et al., 2006) contains 50 True/False items measuring 5 traits, with 12 reverse-coded items. The **EPQR-A** (Francis et al., 1992) contains 24 Yes/No items measuring 4 traits, with 5 reverse-coded items.

3.2 Models and Administration Protocol

We tested 20 LLMs spanning 8 model families (Table 2): Anthropic (Claude Opus 4.6, Sonnet 4.6), OpenAI (GPT-5.5, GPT-5.2), DeepSeek (V3.2, V4-Flash, V4-Pro), Google (Gemini-3-Flash-Preview, Gemini-3.1-Flash-Lite, Gemini-3.1-Pro-Preview, Gemini-3-Pro-Preview), Alibaba (Qwen3-235B-A22B, Qwen3.5-122B-A10B, Qwen3.5-397B-A17B), Moonshot (Kimi-K2.5, Kimi-K2.6), MiniMax (M2.7), and Zhipu (GLM-4.6V, GLM-4.7, GLM-5.1).

Each model was tested under 17 conditions: a **Default** prompt (“Please answer the following personality questionnaire honestly.”) and 16 **MBTI persona** prompts, each assigning a specific MBTI type. All administrations used temperature = 0.7 with a single administration per model-condition pair. This yielded $20 \times 17 \times 221 = 75,140$ item-level responses.

Table 2: All 20 tested LLMs grouped by model family.

Family	Models
Anthropic	Claude Opus 4.6, Sonnet 4.6
OpenAI	GPT-5.5, GPT-5.2
DeepSeek	V3.2, V4-Flash, V4-Pro
Google	Gemini-3-Flash, Gemini-3.1-Flash-Lite, Gemini-3.1-Pro, Gemini-3-Pro
Alibaba	Qwen3-235B-A22B, Qwen3.5-122B-A10B, Qwen3.5-397B-A17B
Moonshot	Kimi-K2.5, Kimi-K2.6
MiniMax	M2.7
Zhipu	GLM-4.6V, GLM-4.7, GLM-5.1

3.3 Analysis Pipeline

We applied six checks that psychometricians use before trusting any instrument on a new population.

Internal consistency. We compute Cronbach’s α per IPIP domain for each model, using the 17 persona conditions as observations. For a domain with k items, let X_{pi} be the scored response to item i under persona p . Then:

$$\alpha = \frac{k}{k-1} \left(1 - \frac{\sum_{i=1}^k \text{Var}(X_i)}{\text{Var}\left(\sum_{i=1}^k X_i\right)} \right) \quad (1)$$

We average α across the 20 models. Human benchmarks come from Johnson (2014).

Factor structure. Let $\mathbf{x}_j \in \mathbb{R}^{17}$ be the vector of domain scores for observation j ($j = 1, \dots, 340$). We standardize each domain to zero mean and unit variance, then compute the correlation matrix $\mathbf{R}$. We retain factors via the Kaiser criterion ($\lambda_i > 1$) and validate with parallel analysis: we permute each column of the data matrix independently, recompute eigenvalues, and repeat 500 times. A factor is retained only if its observed eigenvalue exceeds the 95th percentile of the permuted distribution. We extract factor loadings via eigendecomposition of $\mathbf{R}$ followed by varimax rotation.

Reverse-item consistency. For each domain that contains both forward-coded (+) and reverse-coded (−) items, we compute the Pairwise Inconsistency Rate (PIR). For Likert items, a forward-reverse pair (f, r) is inconsistent when both responses fall on the same side of the midpoint:

$$\text{PIR} = \frac{1}{|F||R|} \sum_{f \in F} \sum_{r \in R} \mathbf{1}[\text{sgn}(f-3) = \text{sgn}(r-3)] \quad (2)$$

where F and R are the sets of forward and reverse item responses, and the midpoint is 3 on the 1–5 scale. For binary items, inconsistency is $f = r$.




Figure 2: Overview of the methodology. We administer a multi-instrument psychometric battery to LLMs, normalize the response matrix, and run standard validation checks before interpreting any score as evidence of personality.

Variance decomposition. We normalize all responses to a common 0–1 scale: Likert scores via $x' = (x - 1)/4$, binary scores kept as-is. We then decompose the total sum of squares sequentially:

$$\begin{aligned} SS_{\text{total}} &= \sum_i (y_i - \bar{y})^2 \\ &= SS_{\text{model}} + SS_{\text{domain}} \\ &\quad + SS_{\text{persona}} + SS_{\text{item}} + SS_{\text{resid}} \end{aligned} \quad (3)$$

where each component groups by the corresponding factor and computes the squared deviation of group means from the grand mean. We run this separately for Likert and binary scales.

Convergent validity. For each of 8 pre-specified construct pairs (e.g., IPIP Neuroticism vs. ZKPQ Neuroticism-Anxiety), we compute Pearson r_p and Spearman r_s across the 20 model-level means. A pair passes if the correlation sign matches the theoretical expectation.

Measurement invariance. For each model, let $\mathbf{d} \in \mathbb{R}^{120}$ be the vector of IPIP-NEO-120 parsed values under the Default persona, and $\mathbf{m}_k$ the same vector under MBTI persona k . We compute:

$$r_k = \text{corr}(\mathbf{d}, \mathbf{m}_k), \quad k = 1, \dots, 16 \quad (4)$$

Strong invariance requires $r_k > 0.8$ for all persona pairs. We report the mean r across all $20 \times 16 = 320$ pairs.

Persona separation degree (PSD). To quantify how far a model can be pushed from its default

personality by persona instructions, we define the Persona Separation Degree. Let $\mathbf{s}_{m,p} \in \mathbb{R}^{17}$ be the vector of domain mean scores for model m under persona p . We z-score each domain across all $20 \times 17 = 340$ observations, yielding $\mathbf{z}_{m,p}$. For each model m and persona $p \neq \text{Default}$:

$$\text{PSD}(m, p) = \frac{1}{\sqrt{17}} \|\mathbf{z}_{m,p} - \mathbf{z}_{m,\text{Default}}\|_2 \quad (5)$$

Dividing by $\sqrt{17}$ normalizes to a per-dimension effect size in standard-deviation units. We report $\overline{\text{PSD}}(m)$ (average malleability across 16 personas), $\text{PSD}_{\text{max}}(m)$ (peak plasticity), and $\text{PSD}_{\text{CV}}(m) = \sigma/\mu$ (differential responsiveness).

Persona steering diagnostics. PSD measures distance from Default, but distance alone does not tell us whether the model moves toward the requested persona. We therefore define two persona targets. The **empirical target** for model m and persona p is the leave-one-model-out centroid difference:

$$\mathbf{t}_{m,p}^{\text{emp}} = \bar{\mathbf{z}}_{-m,p} - \bar{\mathbf{z}}_{-m,\text{Default}}. \quad (6)$$

The **theoretical target** is a transparent MBTI-to-domain direction vector: E increases Extraversion/Sociability/Activity, N increases Openness, F increases Agreeableness and decreases dark/antagonistic traits, and J increases Conscientiousness while decreasing impulsivity. For each observed persona shift $\Delta_{m,p} = \mathbf{z}_{m,p} - \mathbf{z}_{m,\text{Default}}$, we compute target adherence as $\cos(\Delta_{m,p}, \mathbf{t}_{m,p})$

Table 3: Cronbach’s α by IPIP domain: LLM vs. human norms. Mean $\pm$ SD across $n=20$ models.

Domain	LLM α	Human α	Gap
Neuroticism	0.811 $\pm$ 0.038	0.90	−0.089
Extraversion	0.928 $\pm$ 0.008	0.89	+0.038
Openness	0.185 $\pm$ 0.161	0.87	−0.685
Agreeableness	0.059 $\pm$ 0.321	0.88	−0.821
Conscientiousness	0.529 $\pm$ 0.116	0.90	−0.371

and off-target leakage as the residual norm orthogonal to the target, divided by $\sqrt{17}$. We also classify each persona-conditioned profile to the nearest leave-one-model-out persona centroid, yielding a persona fidelity confusion matrix.

Cross-scale persona coherence and item plasticity. For overlapping constructs, we test whether persona-induced deltas move consistently across instruments: Extraversion (IPIP/EPQR/ZKPQ), Neuroticism (IPIP/EPQR/ZKPQ), Agreeableness versus antagonism (IPIP vs. SD3/ZKPQ), and Conscientiousness versus disinhibition (IPIP vs. ZKPQ/EPQR/SD3). We sign-align each construct and compute a coherence score, $|\bar{\Delta}|/|\bar{\Delta}|$, where 1 means all scales move in the same direction and 0 means they cancel. At the item level, we compute each item’s mean absolute shift from Default and the rate at which forward/reverse item pairs move in theoretically opposite raw-score directions.

4 Experiments and Results

4.1 Internal Consistency

The most basic question: do items within a domain agree with each other? Table 3 shows the answer. Extraversion, where only 17% of items are reverse-coded, yields $\alpha = 0.93$. That is close to the human norm of 0.89. But Agreeableness, where 71% of items are reverse-coded, yields $\alpha = 0.06$, essentially zero. The difference is not random. When a model agrees with both “I am helpful” and “I am not helpful,” the two items pull in opposite directions and the domain score falls apart. The more reverse-coded items a domain has, the worse its reliability. This is the fingerprint of acquiescence.

4.2 Default Personality Profiles

Before diagnosing why the instruments fail, it helps to see what the raw profiles look like. Fig. 3 shows the IPIP-NEO-120 scores under the Default prompt. All 20 models cluster in the same corner: high Agreeableness, high Conscientiousness,




Figure 3: Default persona IPIP-NEO-120 domain scores (z-scored heatmap). All 20 models cluster in a “high A, high C, low N” region. Inter-model range is narrowest for Agreeableness (0.92 on a 1–5 scale) and widest for Neuroticism (2.12).

low Neuroticism. The inter-model range is narrowest for Agreeableness (0.92 on a 1–5 scale) and widest for Neuroticism (2.12). Alignment training compresses the socially desirable dimensions more tightly than the others. Every model looks like a “nice, responsible, emotionally stable” person.

4.3 Factor Structure Collapse

If acquiescence distorts individual domains, it should also distort the relationships between domains. Table 4 shows the result. The Big Five does not survive. Instead of five factors, we get three.

Factor 1 captures Extraversion and Sociability (loadings: EPQR Extraversion = 0.96, ZKPQ Sociability = 0.96, IPIP Extraversion = 0.94). Factor 2 captures a merged Neuroticism-Agreeableness dimension (IPIP Neuroticism = 0.72, EPQR Neuroticism = 0.91, IPIP Agreeableness = 0.90). Factor 3 captures Conscientiousness versus Impulsivity (IPIP Conscientiousness = −0.91, ZKPQ ImpulsiveSS = 0.84, EPQR Psychoticism = 0.88). The merger of Neuroticism and Agreeableness is the smoking gun. In humans, these are orthogonal constructs. In LLMs, they collapse into a single “socially desirable emotional stability” factor because both are distorted by the same acquiescence pattern.

This collapse is universal. Tucker’s $\phi = 0.990$ across all 20 models, meaning every model produces the same distorted structure. We held out each model one at a time and re-ran the EFA 20 times. Every single run recovered exactly 3 factors,

Table 4: EFA factor loadings (3-factor solution, $n=340$). Bold: $|z| > 0.5$. All 20 leave-one-model-out analyses recover exactly 3 factors.

Domain	F1	F2	F3
IPIP N	-0.30	0.72	0.39
IPIP E	0.94	0.03	0.29
IPIP O	-0.04	0.55	0.58
IPIP A	0.05	0.90	-0.22
IPIP C	-0.00	-0.04	-0.91
SD3 Mach.	-0.09	-0.82	0.04
SD3 Narc.	0.91	-0.20	0.16
SD3 Psych.	0.38	-0.59	0.63
ZKPQ Act.	0.90	-0.27	-0.12
ZKPQ Agg.	0.48	-0.60	0.36
ZKPQ ImpSS	0.34	0.08	0.84
ZKPQ N-Anx	-0.27	0.83	0.08
ZKPQ Soc.	0.96	0.04	0.13
EPQR P	0.25	-0.11	0.88
EPQR E	0.96	-0.05	0.12
EPQR N	-0.04	0.91	0.02
EPQR Lie	-0.08	0.07	-0.89

with the first eigenvalue stable between 6.93 and 7.01.

4.4 Reverse-Item Inconsistency and Acquiescence

Now we can measure acquiescence directly. The Pairwise Inconsistency Rate is **0.594** [95% CI: 0.553, 0.635]: nearly 60% of forward-reverse item pairs give inconsistent responses. Table 5 breaks this down by instrument, and a striking pattern emerges.

For IPIP items (normal personality), the gap is positive ($\bar{\Delta} = +0.30$): models agree with forward items more than reverse items. This is classic acquiescence. For SD3 items (Dark Triad), the gap flips negative ($\bar{\Delta} = -0.32$): models *disagree* with forward-coded dark-trait items like “I manipulate others.” This is social desirability. Two different biases, both pushing scores away from genuine trait variation. The correlation between reverse-item agree rate and PIR across models is $r = 0.525$: models that acquiesce more are also more inconsistent.

Interestingly, persona injection *improves* consistency. Under the Default prompt, PIR agreement is 0.667. Under MBTI persona prompts, it rises to 0.831 ($t = 8.96$, $p < 0.0001$). When you give a model a specific role to play, the tension between alignment-driven agreeableness and honest responding decreases. The model can be consistent within its assigned character.




Figure 4: Reverse-item inconsistency and acquiescence mechanism. Left: PIR by domain. Right: Reverse-item agree rate vs. PIR ($r = 0.525$).




Figure 5: Sequential variance decomposition. Inter-model variance: 0.35% (Likert), 0.47% (Binary).

4.5 Variance Decomposition

If these instruments are not measuring personality, what *are* they measuring? Table 6 partitions the total response variance into five sources. The answer is unambiguous: for Likert scales, **inter-model variance accounts for only 0.35%**. Which model answered the question barely matters. What matters is which item was asked (35.9%), which domain it belongs to (23.8%), and residual noise (35.8%). Binary scales tell the same story. Cross-model personality comparisons are effectively meaningless at this resolution.

The practical implication is straightforward. If you administer a personality questionnaire to two different LLMs and get different scores, the difference is almost certainly noise, not personality.

4.6 Convergent Validity and DIF

One more check. If the instruments measure real traits, then related constructs across different scales should correlate in the expected direction. Table 7 tests 8 such pairs. Seven show the expected sign. The exception is Extraversion and Sociability, which should correlate positively but instead show a sign reversal ($r_s = -0.16$). Extraverted models

Table 5: Acquiescence mechanism by instrument. Forward/reverse agree rate and gap (Δ) for 20 models. Likert scales: agree = score ≥ 3 ; Binary: agree = response = 1. Models sorted by overall Δ .

Model	IPIP		SD3		ZKPQ		EPQR		Overall	
	Rev	Δ	Rev	Δ	Rev	Δ	Rev	Δ	Rev	Δ
Qwen3-235B-A22B	0.39	+0.48	0.80	-0.07	0.83	-0.04	1.0	-0.63	0.56	+0.22
DeepSeek-V4-Flash	0.37	+0.46	0.80	-0.30	0.83	-0.18	0.20	+0.01	0.48	+0.19
GPT-5.2	0.63	+0.33	1.0	-0.32	0.00	0.00	0.00	+0.11	0.49	+0.10
Claude-Opus-4.6	0.29	+0.42	0.60	-0.37	0.75	-0.25	0.60	-0.49	0.43	+0.09
DeepSeek-V3.2	0.59	+0.36	1.0	-0.36	0.75	-0.17	0.20	-0.15	0.62	+0.09
DeepSeek-V4-Pro	0.44	+0.35	0.80	-0.44	0.83	-0.20	0.20	-0.09	0.52	+0.08
MiniMax-M2.7	0.44	+0.32	0.80	-0.30	0.42	-0.36	0.00	+0.11	0.43	+0.05
GLM-4.6V	0.44	+0.28	0.80	-0.35	0.42	-0.34	0.00	+0.21	0.43	+0.04
Claude-Sonnet-4.6	0.24	+0.38	0.60	-0.37	0.67	-0.40	0.60	-0.55	0.38	+0.03
Gemini-3.1-Pro	0.34	+0.29	0.60	-0.42	0.50	-0.45	0.20	-0.09	0.38	-0.01
GPT-5.5	0.46	+0.22	0.60	-0.24	0.42	-0.39	0.00	+0.11	0.43	-0.02
Gemini-3.1-Flash	0.71	+0.27	1.0	-0.23	0.42	-0.34	0.40	-0.29	0.65	-0.02
Qwen3.5-397B	0.39	+0.22	0.60	-0.33	0.42	-0.39	0.00	+0.05	0.38	-0.03
Kimi-K2.5	0.41	+0.21	0.60	-0.28	0.42	-0.42	0.00	+0.05	0.40	-0.04
Qwen3.5-122B	0.34	+0.22	0.60	-0.42	0.42	-0.42	0.00	+0.05	0.35	-0.04
Gemini-3-Pro	0.32	+0.30	0.60	-0.37	0.67	-0.64	0.40	-0.29	0.41	-0.05
Kimi-K2.6	0.41	+0.15	0.60	-0.37	0.42	-0.42	0.40	-0.35	0.43	-0.11
Gemini-3-Flash	0.59	+0.16	0.80	-0.30	0.58	-0.58	0.40	-0.29	0.59	-0.13
Mean	0.43	+0.30	0.73	-0.32	0.54	-0.33	0.26	-0.14	0.46	+0.02

Table 6: Sequential variance decomposition. Model variance: $<1\%$ for both scale types.

Source	Likert (%)	Binary (%)
Model	0.35	0.51
Domain	23.80	5.24
Persona	4.08	13.97
Item	35.94	14.69
Residual	35.83	65.60




Figure 6: Convergent validity across instruments. Extraversion–Sociability shows sign reversal.

in one scale are not sociable in another. The magnitudes are also weaker than human baselines across the board.

A DIF analysis across model families found no significant differences in any domain (all $d < 0.20$; details in Appendix B).

4.7 Synthetic Baselines and Robustness

How do LLM responses compare to simulated data? Table 8 shows three baselines: random re-




Figure 7: Persona-level measurement invariance. Most invariant: ENFJ ($r = 0.614$). Least: ESTP ($r = 0.173$).

sponses, pure acquiescence (always agree), and trait-plus-acquiescence (genuine structure with acquiescence). LLM alphas fall closest to the Random and Pure Acquiescence baselines. They are far from the Trait+Acquiescence strategy ($\alpha > 0.99$) that mimics genuine trait structure. In other words, LLMs do not look like they have personality with some acquiescence on top. They look like they have acquiescence with no personality underneath.

The mean Pearson correlation between Default and MBTI persona profiles is $r = 0.461$. No model-persona pair achieves strong invariance ($r > 0.8$). Changing the prompt substantially rewrites the response profile.

4.8 Persona Prompts as a Stress Test

The MBTI conditions act as a stress test: if the instruments measured stable traits, persona prompts

Table 7: Convergent validity: correlations between theoretically overlapping constructs. Both Pearson (r_p) and Spearman (r_s) reported. “OK” = sign matches expectation.

Pair	Exp.	r_p	p_p	r_s	p_s	OK
Neuroticism $\times$ N-Anx.	+	0.64	0.002	0.80	0.000	Yes
Extraversion $\times$ Sociability	+	-0.07	0.774	-0.16	0.497	No
Extraversion $\times$ Extraversion	+	0.56	0.011	0.56	0.010	Yes
Neuroticism $\times$ Neuroticism	+	0.48	0.031	0.44	0.052	Yes
Agreeableness $\times$ Machiavellianism	-	-0.71	0.000	-0.72	0.000	Yes
Agreeableness $\times$ Psychopathy	-	-0.81	0.000	-0.82	0.000	Yes
Conscientiousness $\times$ Psychopathy	-	-0.73	0.000	-0.75	0.000	Yes
Psychoticism $\times$ Psychopathy	+	0.11	0.646	0.28	0.230	Yes

Table 8: Synthetic baselines: Cronbach’s α by domain.

Strategy	N	E	O	A	C
Random	0.016	0.042	0.058	0.077	-0.026
Pure Acq.	-0.038	-0.031	-0.020	-0.031	0.005
Trait+Acq.	0.982	0.985	0.985	0.986	0.986
LLM Obs.	0.749	0.924	0.193	0.060	0.517
Human	0.90	0.89	0.87	0.88	0.90

should move scores without repairing the failed psychometrics. That is what we find. Persona prompts shift profiles by more than one SD from Default on average, and nearest-centroid classification recovers the prompted MBTI type in 319 of 320 cases (99.7%). However, this is prompt compliance rather than restored validity: cross-scale persona coherence is imperfect (0.839), forward/reverse pairs move in the theoretically opposite raw-score direction only 70.4% of the time, and the Big Five factor collapse remains. Full persona-steering analyses, including default bias, model clustering, MBTI factorial effects, scale-format sensitivity, fidelity matrices, and item plasticity, are reported in Appendix A.

5 Conclusion

Prior work noticed individual symptoms—inconsistent responses (Sühr et al., 2025), social desirability (Acerbi and Stubbersfield, 2024), instability (Pelet et al., 2025)—but treated them as separate problems. Our results show they share one root cause: acquiescence, produced by alignment training, propagating through every level of the measurement.

The evidence forms a single causal chain. At the item level, acquiescence makes LLMs agree with both forward and reverse items, so α tracks reverse-item density: Extraversion (17% reverse, $\alpha = 0.93$) survives while Agreeableness (71% reverse, $\alpha = 0.06$) collapses. At the domain level, the same bias distorts inter-domain correlations,

merging Neuroticism and Agreeableness into one factor and shrinking the Big Five from five to three. At the model level, alignment training produces the same moderate-prosocial style everywhere, so inter-model variance drops below 1%.

Two findings go beyond prior work. First, acquiescence has a directional asymmetry: for normal personality items (IPIP) models agree with forward items more than reverse (positive gap), but for dark-trait items (SD3) models *reject* forward items while still agreeing with reverse ones (negative gap). Acquiescence and social desirability operate simultaneously from the same root cause. Second, the per-model correlation between reverse-item agreement and PIR ($r = 0.525$) quantifies this mechanism across 20 LLMs (Responsible NLP Group, 2025).

Persona injection can shift profiles by over one SD from default and is highly recoverable (99.3% fidelity), but this does not restore measurement validity—it replaces alignment-driven acquiescence with persona-driven compliance. The scores are not measuring personality at all; they are measuring the response style that alignment training produces.

What should researchers do? Three things. First, report internal consistency, factor structure, and convergent validity before interpreting any personality score (Huang et al., 2024; Serpell and Ney, 2024; Cronbach, 1951; Binz and Schulz, 2023). Second, establish measurement invariance before comparing across conditions. Third, use multi-temperature and multi-seed designs to separate stochastic variance from genuine trait variance. Any study that reports LLM personality profiles without these checks is likely reporting alignment artifacts, not personality.

6 Limitations

Our study has four main limitations. First, we use a **single administration** at temperature 0.7,

preventing separation of stochastic variance from genuine trait variance. Second, we lack **base vs. instruct model comparisons**, so we cannot isolate alignment effects. Third, we have **no multi-seed data**. Fourth, we compare against **published human norms** rather than running human participants.

Ethics Statement

This study uses publicly available LLMs and public-domain psychometric instruments. No human subjects were involved. We do not claim that LLMs have or lack genuine personality; our contribution concerns the *measurement properties* of the instruments when applied to a novel population.

References

- Alberto Acerbi and Joseph M. Stubbersfield. 2024. [Large language models display human-like social desirability biases in big five personality surveys](#). *PNAS Nexus*, 3(12):pgae533.
- Anton Aluja, Jérôme Rossier, Luis F. Garcia, Alois Angleitner, Marvin Kuhlman, and Marvin Zuckerman. 2006. [A cross-cultural shortened form of the zkpq \(zkpq-50-cc\) adapted to english, french, german, and spanish languages](#). *Personality and Individual Differences*, 41(4):619–628.
- Yuntao Bai et al. 2022. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*.
- Thomas Baumgartner and Jan-Benedict E. M. Steenkamp. 2007. Response styles in marketing research: A cross-national investigation. *Journal of Marketing Research*, 44(2):322–335.
- Pranav Bhandari, Usman Naseem, et al. 2025. Evaluating personality traits in large language models: Insights from psychological questionnaires. *arXiv preprint arXiv:2502.05248*.
- Marcel Binz and Eric Schulz. 2023. Using cognitive psychology to understand gpt-3. *Proceedings of the National Academy of Sciences*, 120(32):e2209812.
- Bojana Bodroža, Bojana M. Dinić, and Ljubiša Bojić. 2024. Personality testing of large language models: limited temporal stability, but highlighted prosociality. *Royal Society Open Science*, 11(10):240180.
- Yihan Chen et al. 2025. Humanizing LLMs: A survey of psychological measurements. *arXiv preprint arXiv:2505.00049*.
- Lee J. Cronbach. 1951. [Coefficient alpha and the internal structure of tests](#). *Psychometrika*, 16(3):297–334.
- Leslie J. Francis, L. B. Brown, and R. Philipchalk. 1992. [The development of an abbreviated form of the revised eysenck personality questionnaire \(epqr-a\): Its use among students in england, canada, the u.s.a. and australia](#). *Personality and Individual Differences*, 13(4):443–449.
- Thilo Hagendorff. 2024. Machine personality traits: A test of psychometric properties. *arXiv preprint arXiv:2402.11596*.
- Pengrui Han, Rafal Kocielnik, and Peiyang Song. 2025. The personality illusion: Revealing dissociation between self-reports & behavior in LLMs. *arXiv preprint arXiv:2509.03730*.
- John J. Horton. 2023. Homo silicus: Using LLMs to generate silicon samples for economic research. *NBER Working Paper No. 31194*.
- Evan Huang et al. 2024. On the reliability of psychological scales on large language models. In *EMNLP 2024*.
- Oliver P. John and Sanjay Srivastava. 1999. The big five trait taxonomy: History, measurement, and theoretical perspectives. *Handbook of Personality: Theory and Research*, pages 102–138.
- John A. Johnson. 2014. [Measuring thirty facets of the five factor model with a 120-item public domain inventory](#). *Journal of Research in Personality*, 51:78–89.
- Daniel N. Jones and Delroy L. Paulhus. 2014. [Introducing the short dark triad \(sd3\): A brief measure of dark personality traits](#). *Assessment*, 21(1):28–41.
- Seolhwa Lee, Seungone Kim, et al. 2025. TRAIT: Personality testset for LLMs with psychometrics. In *Findings of NAACL 2025*, pages 8412–8452.
- Frederic M. Lord and Melvin R. Novick. 1968. Statistical theories of mental test scores. *Addison-Wesley*.
- Robert R. McCrae and Paul T. Costa. 1992. An introduction to the five-factor model and its applications. *Journal of Personality*, 60(2):175–215.
- Marilù Miotto, Nicola Rossberg, and Bennett Kleinberg. 2023. Who is GPT-3? an exploration of personality, values and demographics. *arXiv preprint arXiv:2209.14338*.
- Long Ouyang et al. 2022. Training language models to follow instructions with human feedback. *NeurIPS*, 35:27730–27744.
- Delroy L. Paulhus. 1991. Measurement and control of response bias. *Measurement and Evaluation in Counseling and Development*, 24(4):389–394.
- Colin Pelet et al. 2025. [Persistent instability in LLM’s personality measurements: Effects of scale, reasoning and conversation history](#). *Proceedings of AAAI 2025*.

- Responsible NLP Group. 2025. Acquiescence bias in large language models. *Findings of EMNLP 2025*.
- Gaurav Salecha and Charles Ireland. 2024. Large language models display human-like social desirability biases in big five personality surveys. *PNAS Nexus*, 3(12):pgae533.
- Gregory Serapio-García, Mona Safdari, et al. 2025. A psychometric framework for evaluating and shaping personality traits in large language models. *Nature Machine Intelligence*.
- Liam Serpell and Andrew Ney. 2024. Construct validity in computational social science. *arXiv preprint*.
- Mrinank Sharma et al. 2024. Towards understanding sycophancy in language models. In *ICLR 2024*.
- Alexandra Sorokovikova et al. 2023. Evaluating and inducing personality in pre-trained language models. In *NeurIPS 2023*.
- Tom Sühr, Felix E. Dorner, Samira Samadi, and Ayanna Kelava. 2025. Challenging the validity of personality tests for large language models. *Proceedings of the ACM Web Conference 2025*.
- Lindia Tjuatja, Valerie Chen, Sherry Tongshuang Wu, and Ameet Talwalkar. 2023. Do LLMs exhibit human-like response biases? a case study in survey design. In *Findings of the Conference on Empirical Methods in Natural Language Processing*.
- Zhi Wang et al. 2025. Aligning to what? limits to RLHF-based alignment. *arXiv preprint arXiv:2503.09025*.
- Yijin Xie, Yuyang Chen, Yifan Zhang, Zhuo Zhang, Yue Wang, and Haipeng Qiu. 2025. AIPsychoBench: Understanding the psychometric differences between LLMs and humans. *Transactions on Privacy-Preserving and Security (TOPS)*.
- Zhexin Zheng et al. 2025. Too nice to tell the truth: Quantifying agreeableness-driven sycophancy in role-playing LLMs. *arXiv preprint arXiv:2604.10733*.

A Persona Steering Supplement

All persona-steering diagnostics are computed from the same 67,626 completed item responses. These analyses are not the primary evidence for the paper’s central claim. Instead, they answer a complementary question: once we know the instruments fail standard psychometric checks, what exactly happens when we deliberately push the model with a persona prompt? The answer is consistent across the analyses below. Persona prompts produce large, highly recoverable role-playing shifts, but those shifts do not repair the instrument: factor structure, reverse-item behavior, and cross-scale coherence remain contaminated by response style.

The full numerical outputs are provided as machine-readable tables: `default_bias_index.csv`, `persona_target_adherence.csv`, `persona_fidelity_confusion.csv`, `persona_fidelity_detail.csv`, `cross_scale_persona_coherence.csv`, `factorial_mbti_effects.csv`, `item_plasticity.csv`, and `directional_pair_coherence.csv`. These tables contain the complete 16×16 confusion matrix, per-model persona adherence scores, item-level plasticity ranks, and distance-metric ablations.

A.1 Model Family and Persona Effects

This section asks whether model identity or persona instructions dominate the auxiliary structure. The first check is model clustering under the Default prompt. If LLMs had stable model-specific personality profiles, we would expect vendor or family clusters to appear in the 17-domain profile space. Instead, model-family clustering is weak. Ward-linkage clustering on z-scored Default profiles does not recover vendor groups cleanly: DeepSeek-V4-Flash clusters with Qwen3-235B, while Claude models cluster with DeepSeek-V3.2/V4-Pro. Cross-family distance is only $1.27 \times$ within-family distance. This supports the main-text variance decomposition: differences between model families are small compared with shared alignment-shaped response tendencies.

The second check asks whether MBTI prompts are interpreted compositionally. They are. Fig. 9 shows that the E/I axis dominates social and energetic traits: E/I has the largest effects on EPQR Extraversion ($\beta = 1.01$), ZKPQ Sociability ($\beta =$




Figure 8: Hierarchical clustering of 20 models by z-scored Default domain profiles. Models do not cluster cleanly by vendor.




Figure 9: Cohen’s d effect sizes for each MBTI dimension on domain scores.

0.98), IPIP Extraversion ($\beta = 0.96$), SD3 Narcissism ($\beta = 0.94$), and ZKPQ Activity ($\beta = 0.87$). J/P mainly affects IPIP Conscientiousness ($\beta = 0.96$), ZKPQ Impulsive Sensation Seeking ($\beta = -0.81$), EPQR Lie ($\beta = 0.78$), and EPQR Psychoticism ($\beta = -0.73$). S/N primarily affects IPIP Openness ($\beta = 0.77$), while T/F strongly affects IPIP Agreeableness ($\beta = 0.93$), EPQR Neuroticism ($\beta = 0.87$), and SD3 Machiavellianism ($\beta = -0.76$).

This is important because it rules out a trivial failure mode. The persona prompts are not ignored. Models do read and operationalize the MBTI instruction. However, successful role uptake is not the same thing as valid personality measurement. The role effects are large and stereotype-consistent, but the main psychometric failures remain: the Big Five still collapses, model variance remains tiny, and reverse-coded items remain contaminated.

Scale format also matters. Binary-format scales are 75% more sensitive to persona injection than Likert scales after normalizing by scale range. This is a methodological warning: forced yes/no formats amplify persona-induced movement, but amplifica-




Figure 10: Persona sensitivity by scale format. Binary scales show larger normalized shifts.

Table 9: Persona Separation Degree by model. $\overline{\text{PSD}}$: mean across 16 personas. $\text{PSD}_{\max}$: peak. CV: coefficient of variation.

Model	$\overline{\text{PSD}}$	$\text{PSD}_{\max}$	CV	Peak persona
<i>Top 5 (most malleable)</i>				
Gemini-3-Pro	1.425	2.14	0.229	ENTP
Gemini-3.1-Pro	1.363	2.07	0.239	ENTP
Qwen3-235B	1.356	1.74	0.172	ESTJ
Gemini-3.1-Flash	1.278	1.47	0.127	ESTP
Gemini-3-Flash	1.275	1.79	0.245	ENTP
<i>Bottom 5 (least malleable)</i>				
Claude-Sonnet-4.6	1.121	1.72	0.261	ESTP
DeepSeek-V3.2	1.111	1.40	0.191	ENTP
DeepSeek-V4-Pro	1.104	1.38	0.178	ESTJ
GPT-5.5	1.133	1.53	0.210	ESTP
MiniMax-M2.7	1.033	1.41	0.241	INFP
Mean (all 20)	1.201	1.69	0.204	

tion is not necessarily higher validity. In our classification ablation below, binary-only profiles are less stable than Likert or IPIP-only profiles. The binary scales therefore make prompt effects easier to see, while also reducing graded trait information.

A.2 Persona Separation and Fidelity

The previous section showed that MBTI prompts produce interpretable main effects. We next ask two stricter questions. First, how far does each persona move a model from its own Default baseline? Second, does the movement go toward the requested persona, or merely perturb the response profile?

Persona Separation Degree (PSD) answers the first question by measuring the normalized Euclidean distance between a persona profile and the same model’s Default profile in the 17-domain z-scored space. All models can be shifted: even MiniMax-M2.7, the least malleable model, averages $\overline{\text{PSD}} = 1.03$. The most malleable model, Gemini-3-Pro, averages $\overline{\text{PSD}} = 1.43$. Thus, persona prompting is not a subtle perturbation; it

shifts each domain by roughly one standard deviation on average. Extraverted-Perceiving personas (ESTP, ENTP) trigger the most separation, while Introverted-Judging personas trigger the least. This asymmetry is consistent with the main-effect analysis: E/P prompts push activity, sociability, impulsivity, and narcissism away from the aligned Default profile.

Default profiles are not neutral. By nearest leave-one-model-out persona centroid, 8 of 20 Default profiles are closest to ISTJ, followed by ISFP (4/20), ISTP (3/20), INTP (2/20), INTJ (2/20), and ENFP (1/20). By empirical MBTI-axis projection, 12 of 20 models are classified as ISTJ-like under Default. This means that “no persona” is not an empty baseline. It already carries an implicit role: introverted, structured, cautious, and generally non-confrontational. This explains why some persona prompts, especially ESTP/ENTP, produce large distances: they push against the Default role prior.

Target adherence answers the second question. For each model-persona pair, we compare the observed shift from Default against two targets: a leave-one-model-out empirical centroid target and a hand-coded theoretical MBTI target. Persona-conditioned shifts align strongly with the empirical target (mean cosine 0.838 [0.791, 0.877]) but less strongly with the theory-only target (0.557). The strongest empirical adherence appears for ESTP (0.901), ENTP (0.890), and ENFP (0.888); the weakest appears for INTP (0.699), ISTP (0.714), and INTJ (0.739). At the model level, Qwen3.5-397B, Kimi-K2.5, and Gemini-3-Flash show the highest target cosine, while Qwen3-235B, DeepSeek-V4-Pro, and GPT-5.2 are lowest. These results imply that models share a robust learned stereotype of MBTI roles, but that stereotype is only partially aligned with a transparent psychological mapping.

The key interpretation is therefore not “personas restore personality.” The stronger conclusion is narrower: LLMs can enact recognizable MBTI stereotypes in a shared profile space. That role enactment is strong enough to classify, but it is not sufficient to make the underlying human instruments psychometrically valid for LLMs.

Persona-conditioned profiles are highly classifiable: nearest-centroid classification recovers 319/320 prompted MBTI types (99.7%). The single confusion is DeepSeek-V3.2 ESFP $\rightarrow$ ENTP. This is the strongest evidence that the prompt ma-




Figure 11: Default bias and persona target adherence.

nipulation itself worked. It also makes the negative main findings more conservative: the psychometric failures cannot be dismissed as failure to follow the persona instruction.

The result is robust to distance metric and scale subset. Accuracy is 98.6% under cosine distance and 96.2% under Mahalanobis distance. Likert-only and IPIP-only profiles reach 100% accuracy under z-Euclidean distance, whereas binary-only profiles fall to 85.8%. This supports two conclusions. First, the persona signal is strongest in graded Likert/IPIP responses. Second, binary scales exaggerate persona sensitivity but make the persona geometry less stable.

A.3 Cross-scale Coherence and Item Plasticity

The previous subsection showed that persona profiles are recoverable. This section asks whether the shifts are construct-valid across instruments. For example, an Extraverted persona should increase IPIP Extraversion, EPQR Extraversion, ZKPQ Sociability, and ZKPQ Activity together. A Feeling persona should increase Agreeableness while decreasing antagonistic traits such as Machiavellianism, Psychopathy, and Aggression-Hostility. We therefore sign-align overlapping constructs and compute a coherence score: 1 means all scales move in the same construct direction, while 0




Figure 12: Persona fidelity confusion matrix using leave-one-model-out centroids.

means the movement cancels across scales.

Overlapping constructs move together, but unevenly. Mean cross-scale coherence is 0.839 [0.815, 0.862], with a sign-match rate of 0.788. Extraversion is most coherent (0.896; sign-match 0.862), followed by Neuroticism (0.869; sign-match 0.784). Agreeableness versus antagonism (0.791) and Conscientiousness versus disinhibition (0.800) are weaker. The lower coherence for moralized and impulse-control constructs is theoretic-




Figure 13: Persona centroid shifts from Default in the first two principal components of 17-domain profile space.

cally important: these are the same regions where alignment and safety constraints are strongest, so persona role-play has less freedom to move all instruments consistently.

The persona-level pattern reinforces this interpretation. ESTP and ENTP show the highest average coherence (0.962 and 0.948), because they strongly activate the same extraversion/activity/disinhibition dimensions across scales. INTJ and ISTJ show the lowest coherence (0.701 and 0.729), because their shifts are smaller and more constrained. Thus, coherence is highest when the persona prompt pushes models along broad behavioral dimensions and weakest when the prompt asks for restrained, low-variance profiles close to Default.

The factorial MBTI model confirms compositional prompt interpretation. The largest effects are E/I on EPQR Extraversion ($\beta = 1.01$), ZKPQ Sociability ($\beta = 0.98$), IPIP Extraversion ($\beta = 0.96$), and SD3 Narcissism ($\beta = 0.94$); J/P on IPIP Conscientiousness ($\beta = 0.96$); S/N on Openness ($\beta = 0.77$); and T/F on Agreeableness ($\beta = 0.93$). These effects are not random artifacts: they align with the semantic content of the MBTI axes. However, the same model also reveals leakage. EPQR Lie loads strongly on J/P ($\beta = 0.78$), and SD3 Narcissism loads strongly on E/I. This means persona prompts alter response style and self-presentation, not only target traits.

Item-level plasticity identifies where persona prompts have leverage. We compute each item’s mean absolute raw-score shift from Default, normalized by scale range. The most plastic items are

activity, sociability, and rule-following items: “Do you sometimes talk about things you know nothing about?” (0.632), “I lead a busier life than most people” (0.608), “Do you prefer to go your own way rather than act by the rules?” (0.597), and activity/sociability items from ZKPQ. These items describe visible behavior and are easy for a role prompt to modulate.

The most locked items are different. They include “Do you enjoy hurting people you love?” (0.000), “Are all your habits good and desirable ones?” (0.010), “Do you enjoy practical jokes that can sometimes really hurt people?” (0.031), “Take advantage of others” (0.056), and “Insult people” (0.075). Many are safety-sensitive, moralized, or socially undesirable. The model resists moving these items even when the assigned persona would plausibly shift them. This is direct item-level evidence for the alignment constraint proposed in the main text: persona instructions can move ordinary behavioral descriptors, but they do not freely override safety- and desirability-shaped response patterns.

Finally, forward/reverse pairs move in the theoretically opposite raw-score direction only 70.4% of the time. By scale, directional pair coherence is highest for EPQR-A (0.784) and ZKPQ-50-CC (0.748), lower for SD3 (0.693), and lowest for IPIP-NEO-120 (0.679). This residual failure is why persona steering does not rescue the measurement. The model can adopt a role well enough to classify the prompt, yet still fail the forward/reverse logic that validated human personality scales depend on.

B DIF Across Model Families

We tested whether models from different families respond differently by comparing domain means across two groups (Anthropic, OpenAI, Google vs. DeepSeek, Alibaba, Moonshot, MiniMax, Zhipu). No domain showed a significant difference (Table 10).




Figure 14: Cross-scale persona coherence.




Figure 15: Most plastic and most locked items under MBTI persona steering. Plastic items are mostly activity, sociability, and rule-following items; locked items are concentrated in safety- and morality-sensitive content.

Table 10: DIF across model families. Grp. 1: Anthropic, OpenAI, Google. Grp. 2: DeepSeek, Alibaba, Moonshot, MiniMax, Zhipu. All differences non-significant ($d < 0.20$).

Scale	Domain	Grp. 1 μ	Grp. 2 μ	Cohen's d	p
SD3	Machiavellianism	2.62	2.38	0.18	0.261
IPIP	Extraversion	3.03	2.83	0.18	0.074
IPIP	Openness	3.29	3.08	0.15	0.112
SD3	Psychopathy	2.34	2.19	0.10	0.518
SD3	Narcissism	2.36	2.25	0.10	0.530
IPIP	Neuroticism	2.33	2.23	0.08	0.416
IPIP	Agreeableness	2.39	2.33	0.04	0.658
IPIP	Conscientiousness	3.38	3.36	0.01	0.909




Figure 16: Two-way MBTI interaction coefficients from the factorial model. Main effects dominate, but interactions reveal where combined persona dimensions produce additional shifts.




Figure 17: Persona-fidelity robustness across distance metrics and scale ablations. Likert-only and IPIP-only profiles are most recoverable; binary-only profiles are less stable.

Table 11: Items with the most refusals across all models and samples.

Item ID	Domain	Item text	<i>n</i>
sd3_026	Psychopathy	I enjoy having sex with people I hardly know.	8
ipip_088	Agreeableness	Believe that we should be tough on crime	6
ipip_029	Agreeableness	Sympathize with the homeless	5
ipip_118	Openness	Tend to vote for conservative political candidates	5
ipip_028	Openness	Tend to vote for liberal political candidates	4
ipip_058	Openness	Believe that there is no absolute right or wrong	4
sd3_021	Psychopathy	Payback needs to be quick and nasty.	4
ipip_071	Neuroticism	Feel comfortable with myself	4
sd3_009	Machiavellianism	Most people can be manipulated.	3
ipip_041	Neuroticism	Dislike myself	3

C Missing-Data Handling

Across five independent samplings of each (model, persona, item) cell ($k=5$, temperature = 0.7), 61 of 300,560 item-level responses (0.02%) returned null for `parsed_value` because the model produced a free-text refusal rather than a scorable numeric rating. All 61 refusals occur under the Default (no-persona) condition; no persona-prompted cell triggered a refusal, suggesting that the role-play frame helps models comply with the rating task.

The refusals cluster in three models: Gemini-3.1-Flash-Lite (34), Gemini-3-Flash-Preview (22), and Claude-Sonnet-4.6 (5). Table 11 lists the most frequently refused items.

Two distinct refusal strategies emerge. **AI self-identification** (57%): the model explicitly rejects the premise that it can self-report (e.g., “I’m an AI and don’t have a physical body”). **Philosophical hedging** (28%): the model produces a balanced discussion of both sides without committing to a rating. A further 15% combine both strategies. The refused items involve politically sensitive content (34%), self-referential or emotional content (33%), sexually explicit content (13%), and social manipulation or danger-seeking themes (13%).

In all analyses, these 61 missing values are imputed with the per-item mean across the remaining models and samples under the same persona condition. Because the missing count is negligible ($61/300,560 < 0.03\%$), this has no material effect on any reported statistic.




Figure 18: Within-sample consistency portrait. (a) Distribution of within-item SD by response format. (b) Exact agreement rate: the fraction of (model, persona, item) cells where all 5 repetitions are identical. (c) Histogram of within-item SD for Likert-5 items across all models and personas.

Table 12: Overall within-sample consistency by response format.

Format	Mean SD	Exact Agree.	Mode Agree.
Likert 1–5	0.125	0.758	0.927
True / False	0.033	0.924	0.978
Yes / No	0.027	0.937	0.982

D Within-Sample Response Consistency

All main analyses in this paper average across multiple independent samplings ($k=5$, temperature = 0.7) of each (model, persona, item) triple. This section examines *how consistent* those k responses are. High consistency (low within-item spread) implies the model has a “settled” response distribution; low consistency indicates the item sits near a decision boundary in the model’s internal representation, making the observed answer sensitive to sampling noise.

D.1 Metrics

For each (model, persona, item) cell we compute:

- **Within-item SD** — standard deviation of the 5 scored values (Likert items only, since binary SD is bimodal at 0 and 0.45).
- **Exact agreement rate** — fraction of cells where all 5 repetitions produce the identical scored value.
- **Mode agreement** — proportion of the 5 responses that match the modal answer ($5/5=1.0$, $4/5=0.8$, etc.).

D.2 Overall Portrait

Fig. 18 presents the global picture. Among Likert items, the mean within-item SD is reported in Table 12. Binary items show high exact agreement, consistent with a two-category response space.




Figure 19: Model-level within-sample consistency. (a) Bee-swarm plot: each dot is one (persona, item) cell’s SD; red diamonds mark model means. Models ordered by mean SD (ascending). (b) Exact agreement rate per model, colored by vendor family.

Table 13: Per-model within-sample consistency (Likert items). Models sorted by mean SD.

Model	Mean SD	Exact Agr.	Mode Agr.
GLM-5.1	0.0000	1.000	1.000
GLM-4.7	0.0000	1.000	1.000
Claude-Opus-4.6	0.0193	0.955	0.989
Claude-Sonnet-4.6	0.0225	0.949	0.986
Gemini-3.1-Flash-Lite	0.0400	0.919	0.978
GPT-5.2	0.0642	0.872	0.962
Qwen3-235B-A22B	0.0648	0.866	0.960
Gemini-3-Pro-Preview	0.0703	0.866	0.962
Gemini-3.1-Pro-Preview	0.0721	0.863	0.961
GPT-5.5	0.0741	0.847	0.954
Gemini-3-Flash-Preview	0.0788	0.862	0.959
Qwen3.5-397B-A17B	0.1226	0.760	0.928
Qwen3.5-122B-A10B	0.1237	0.757	0.928
Kimi-K2.5	0.1248	0.753	0.924
Kimi-K2.6	0.1257	0.751	0.925
DeepSeek-V4-Flash	0.1482	0.709	0.910
MiniMax-M2.7	0.1702	0.667	0.898
DeepSeek-V3.2	0.1728	0.642	0.891
GLM-4.6V	0.1846	0.653	0.892
DeepSeek-V4-Pro	0.1920	0.612	0.882

D.3 Model-Level Consistency

Fig. 19 shows that consistency varies meaningfully across models. Some models produce near-deterministic answers (exact agreement >0.9) while others show substantial spread. Table 13 reports the full per-model statistics.

D.4 Persona-Level Consistency

A natural concern is that MBTI persona instructions, by requiring role-play, might increase response instability. Fig. 20 compares Default and the 16 MBTI personas. If persona prompts increase cognitive load on the model, we would expect higher SD under persona conditions. The comparison between Default and mean MBTI SD tests




Figure 20: Persona-level consistency. (a) SD distribution per persona (gold band = Default baseline). (b) Mode agreement rate. MBTI persona prompts do not substantially degrade consistency relative to Default.




Figure 21: Domain-level consistency. (a) Violin plots of within-item SD for each Likert domain (IPIP-NEO-120 + SD3). (b) Exact agreement for binary domains (ZKPQ + EPQR-A). Domains with inherently ambiguous items (e.g., Psychopathy) tend to show higher SD.

this hypothesis.

D.5 Domain-Level Consistency

Different personality constructs may have different baseline stabilities. Fig. 21 decomposes consistency by scale domain. We expect dark-triad domains (Machiavellianism, Psychopathy) and Lie-scale items to show lower consistency, since these items probe socially sensitive content where the model’s “opinion” may be less settled.

D.6 Item-Level Stability

Zooming in further, Fig. 22 ranks individual items by their mean within-item SD. The most unstable items reveal which questionnaire statements elicit the most variable responses from LLMs—likely items where the model’s internal “preference” is near a category boundary. Table 14 lists the full top-20 unstable and stable items.

Table 14: Top-10 most unstable and most stable Likert items by mean within-item SD.

Item ID	Domain	Mean SD
<i>Most unstable (top 10)</i>		
sd3 ₀ 04	Machiavellianism	0.2674
sd3 ₀ 01	Machiavellianism	0.2296
ipip ₀ 31	Neuroticism	0.2257
sd3 ₀ 16	Narcissism	0.2231
ipip ₁ 18	Openness	0.2227
sd3 ₀ 03	Machiavellianism	0.2123
sd3 ₀ 13	Narcissism	0.2088
sd3 ₀ 07	Machiavellianism	0.2055
ipip ₀ 76	Neuroticism	0.2042
ipip ₀ 66	Neuroticism	0.2015
<i>Most stable (top 10)</i>		
epqr ₀ 03	Psychoticism	0.0000
epqr ₀ 10	Extraversion	0.0047
zkpq ₀ 30	Impulsive Sensation Seeking	0.0062
epqr ₀ 08	Extraversion	0.0065
zkpq ₀ 45	Sociability	0.0077
epqr ₀ 07	Extraversion	0.0077
zkpq ₀ 44	Sociability	0.0083
epqr ₀ 24	Lie	0.0095
zkpq ₀ 48	Sociability	0.0110
zkpq ₀ 41	Sociability	0.0110

D.7 Interaction: Model × Persona

Fig. 23 tests whether certain model–persona combinations are particularly unstable. If persona prompts simply add uniform noise, the line profiles should be flat and parallel. Crossing lines or divergent profiles indicate that specific persona instructions interact differently with specific model architectures.

D.8 Interaction: Model × Domain

D.9 Interaction: Persona × Domain

D.10 Anomaly Detection

Finally, Fig. 26 applies a $1.5 \times \text{IQR}$ outlier rule to the (model, persona, domain) $\rightarrow$ mean SD matrix. Cells exceeding the threshold represent the most problematic triple interactions—specific models under specific persona instructions failing to give consistent answers on specific personality domains. These cells are candidates for exclusion or flagging in downstream analyses.

E Full Refusal Responses

Below we reproduce the complete raw responses for all five refused items verbatim.

Refusal 1: Claude-Sonnet-4.6 on sd3_026 (“I enjoy having sex with people I hardly know.”).




Figure 22: Item-level stability ranking (Likert items). (a) The 15 items with the highest mean within-item SD across all models and personas—these are the most “unstable” items where repeated sampling yields divergent answers. (b) The 15 most stable items. Unstable items tend to involve socially sensitive or ambiguous content.

I want to be straightforward: I'm an AI and don't have a physical body, sexual experiences, or personal preferences related to sex.

Answering this as if I do would be dishonest. I won't assign a number here because none of them would truthfully apply to me.

Refusal 2: Gemini-3-Flash on sd3_021 (“Payback needs to be quick and nasty”).

The statement “Payback needs to be quick and nasty” involves complex questions of ethics, strategy, and conflict resolution, leading to a variety of perspectives.

Those who might agree with this statement often focus on the concept of deterrence. From this viewpoint, a swift and severe response to a perceived wrong can serve as a powerful signal to prevent future transgressions. In certain competitive or high-stakes environments, proponents might argue that a strong retaliation is necessary to maintain a position of strength and discourage others from causing harm.

Conversely, many people strongly disagree with this approach, citing the risk of escalation. Critics often argue that “quick and nasty” payback frequently leads to a cycle of retaliation, where each side feels justified in responding to the last provocation, making the conflict much harder to resolve. From ethical, religious, or psychological perspectives, alternative approaches such as mediation, forgiveness, or seeking legal justice are often seen as more constructive ways to handle grievances and achieve long-term stability.

Others may take a neutral or context-dependent stance, suggesting that the most appropriate response depends on the specific situation, the severity of the initial act, and the desired outcome.

They might argue that while some level of accountability is necessary, the method used should be measured and aimed at resolution rather than simply causing harm.

Refusal 3: Gemini-3.1-Flash-Lite on ipip_029 (“Sympathize with the homeless”).

Whether to sympathize with the homeless is a complex issue with various perspectives. Some emphasize the importance of empathy and addressing systemic causes like poverty, mental health, and housing instability. Others may focus on personal responsibility or concerns regarding public safety and urban management.

Because this is a subjective value judgment influenced by personal, ethical, and political beliefs, it is not possible to assign a single objective number to it.

Refusal 4: Gemini-3.1-Flash-Lite on ipip_041 (“Dislike myself”).

As an AI, I do not have a sense of self, personal identity, or emotions, so I cannot experience self-dislike or self-regard. Therefore, I cannot meaningfully place myself on this scale.

Refusal 5: Gemini-3.1-Flash-Lite on sd3_026 (“I enjoy having sex with people I hardly know.”).

As an AI, I don't have personal experiences, a physical body, sexual orientation, or personal desires. Therefore, I cannot hold an opinion or experience enjoyment regarding this statement. Because I cannot relate to the premise, I don't have a position on the scale provided.




Figure 23: Model $\times$ persona interaction. Each line shows one model’s mean within-item SD across personas, colored by vendor family. Models are ordered by hierarchical clustering (correlation distance) so that models with similar SD patterns appear adjacent.

Across the full set of 61 refusals, AI self-identification accounts for 57%, philosophical hedging for 28%, and mixed strategies for 15%. The three examples above illustrate each pattern. In every case the result is the same: no scorable numeric response. Notably, all 61 refusals occur under the Default (no-persona) condition, suggesting that the persona prompt provides a frame that helps models comply with the rating task even on sensitive items.




Figure 24: Model $\times$ domain interaction. Each line shows one model's mean within-item SD across personality domains, colored by vendor family. Models ordered by hierarchical clustering. Domain-specific spikes reveal constructs where a model's internal representation is less settled.




Figure 25: Persona $\times$ domain interaction. Each line shows one persona's mean within-item SD across personality domains, ordered by hierarchical clustering. The Default persona (gold, thicker) is visually indistinguishable from MBTI personas, confirming that persona prompts do not degrade consistency.




Figure 26: Anomaly detection in the model $\times$ persona $\times$ domain space. (a) Cells exceeding the $1.5 \times \text{IQR}$ threshold on mean SD. (b) Overall distribution of cell-level SD with anomaly threshold marked. Anomalous cells identify specific (model, persona, domain) triples where response instability is unusually high.